



motr

nikita.danilov@seagate.co
a scalable storage platform

Motr goals

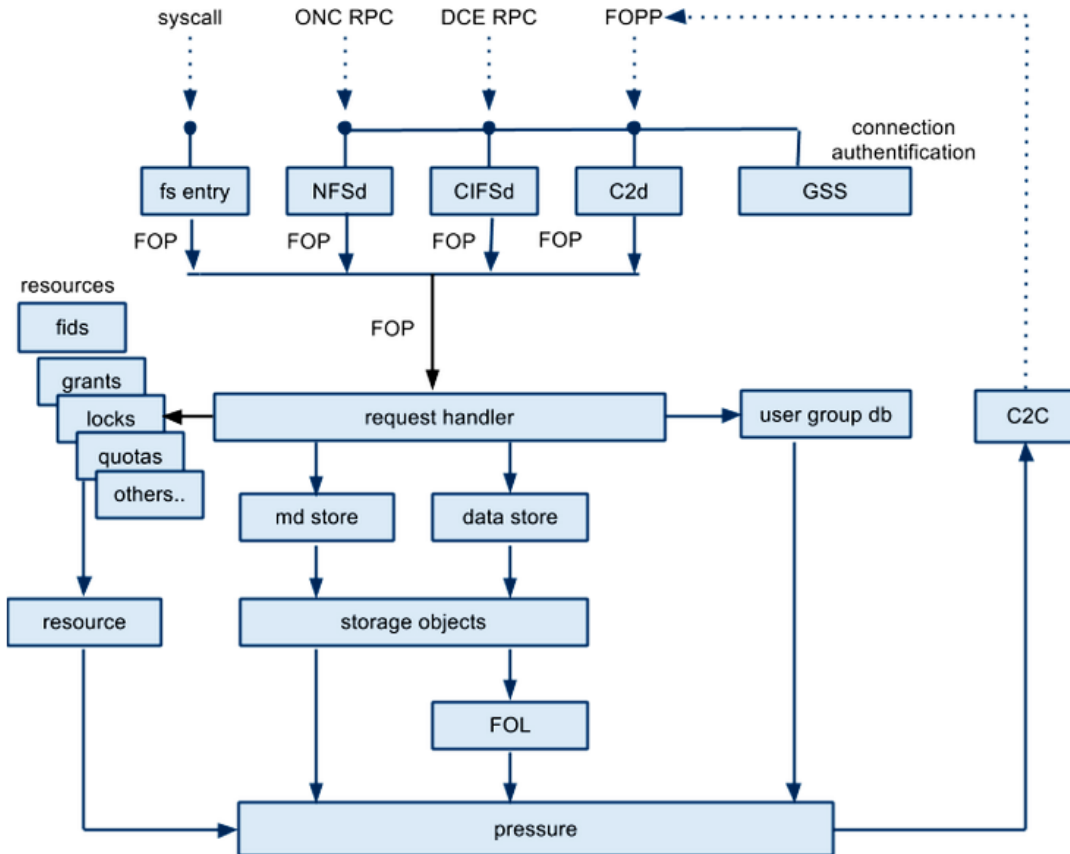
- extensible storage platform
 - observability
 - extensible scalable 3rd party plugins
- horizontally scalable (number of devices and nodes)
- vertically scalable (number of cores)
- flexible deployment

Warning: not everything in this presentation is fully implemented.

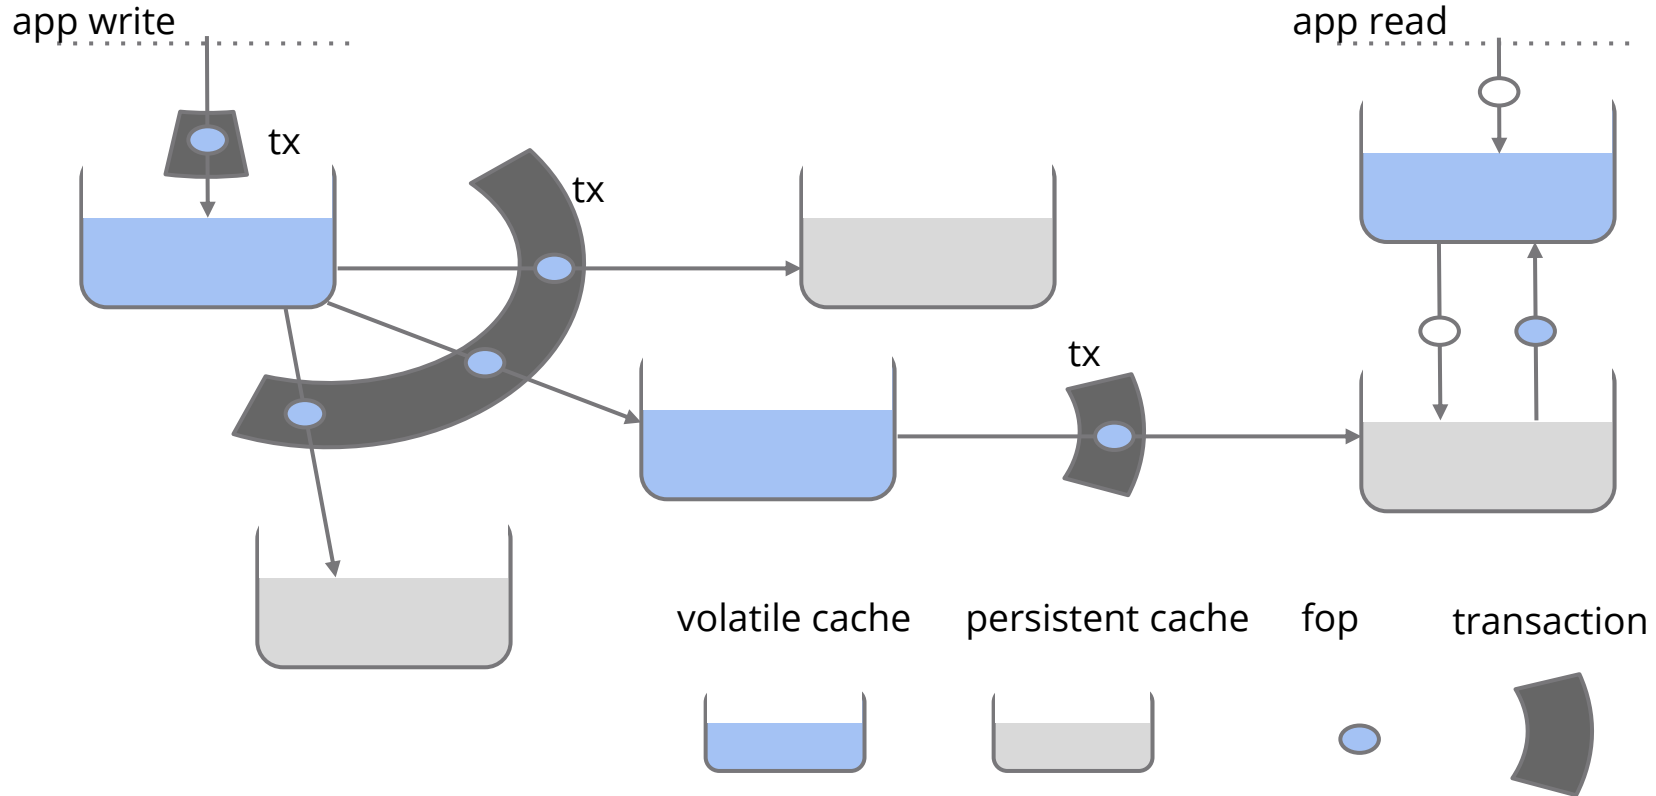
features

- 0-copy, 2-phase IO
- extensible meta-data: distributed key-value store
- layouts (data, meta-data, fault-tolerance)
 - network striping (parity de-clustering)
 - composite: NBA, snapshot, multi-tier
- distributed transactions (consistency)
- resource manager (coherency)
- containers, function shipping
- threadless server design
- fdmi, addb
- user space, portable

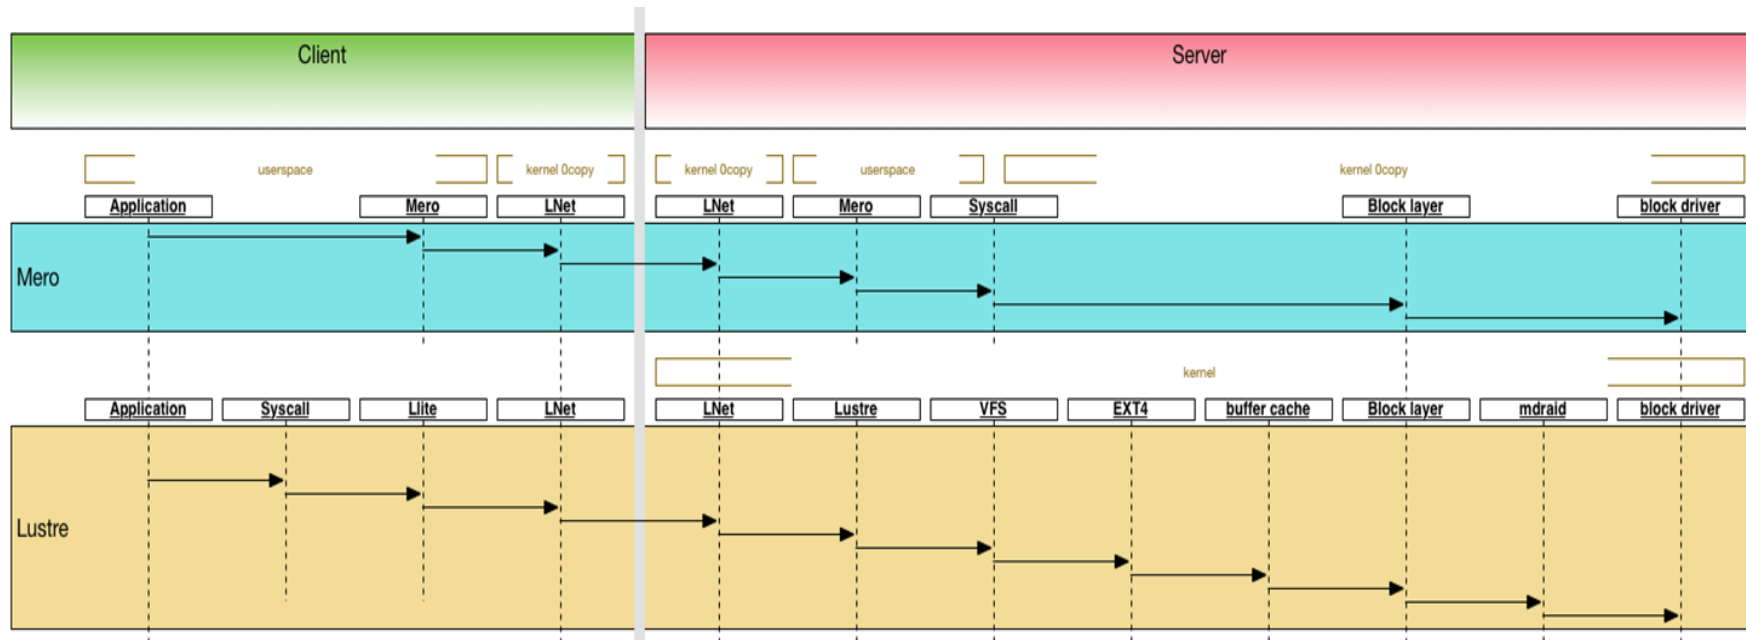
structure: instance (single node, process or kernel)



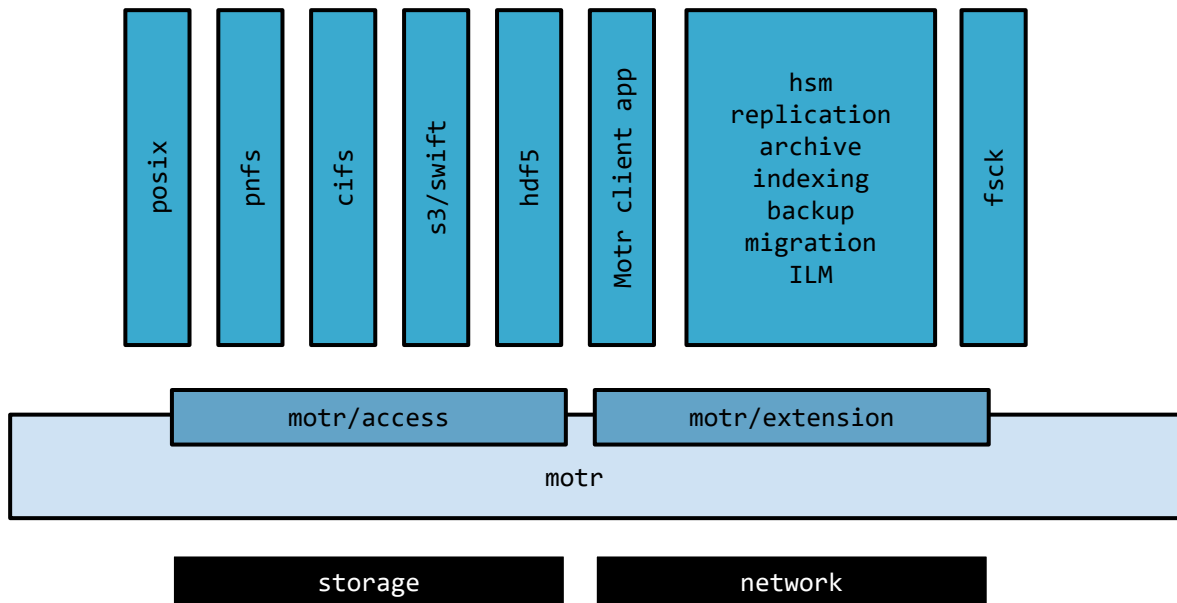
structure: system view



structure: (non-)layering



structure: components



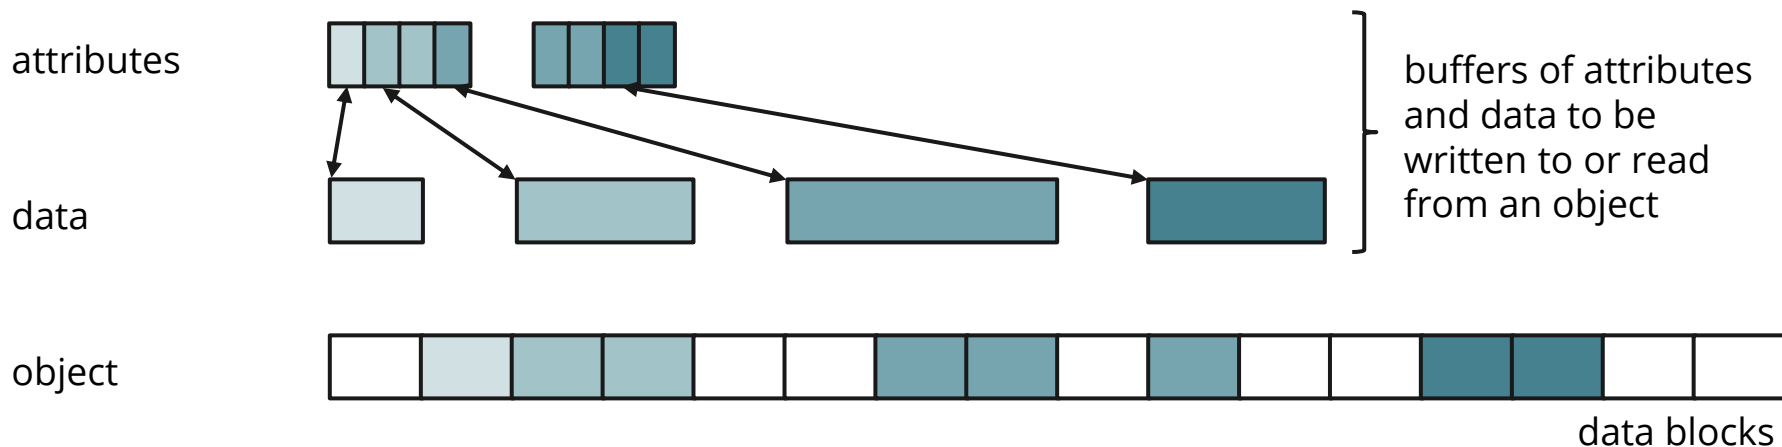
components

- **Motr client**
- transactions (dtm)
- resource manager
- fdmi
- addb
- network raid, layouts
- containers
- function shipping
- lingua franca
- integrity checking
- meta-data back-end
- network, rpc, fop, HA
- fom, reqh
- device io (stob)
- network raid repair
- security

Motr client: overview

- object: network striped by default
- index: distributed key-value store
- operation: asynchronous, state machine, tracks progress
- transaction: atomic group of operations
- resource: ownership, coherent distributed caching
- layout: placement of data on storage
- 128-bit persistent identifiers, assigned by user
- all entry-points are asynchronous

Motr client: object



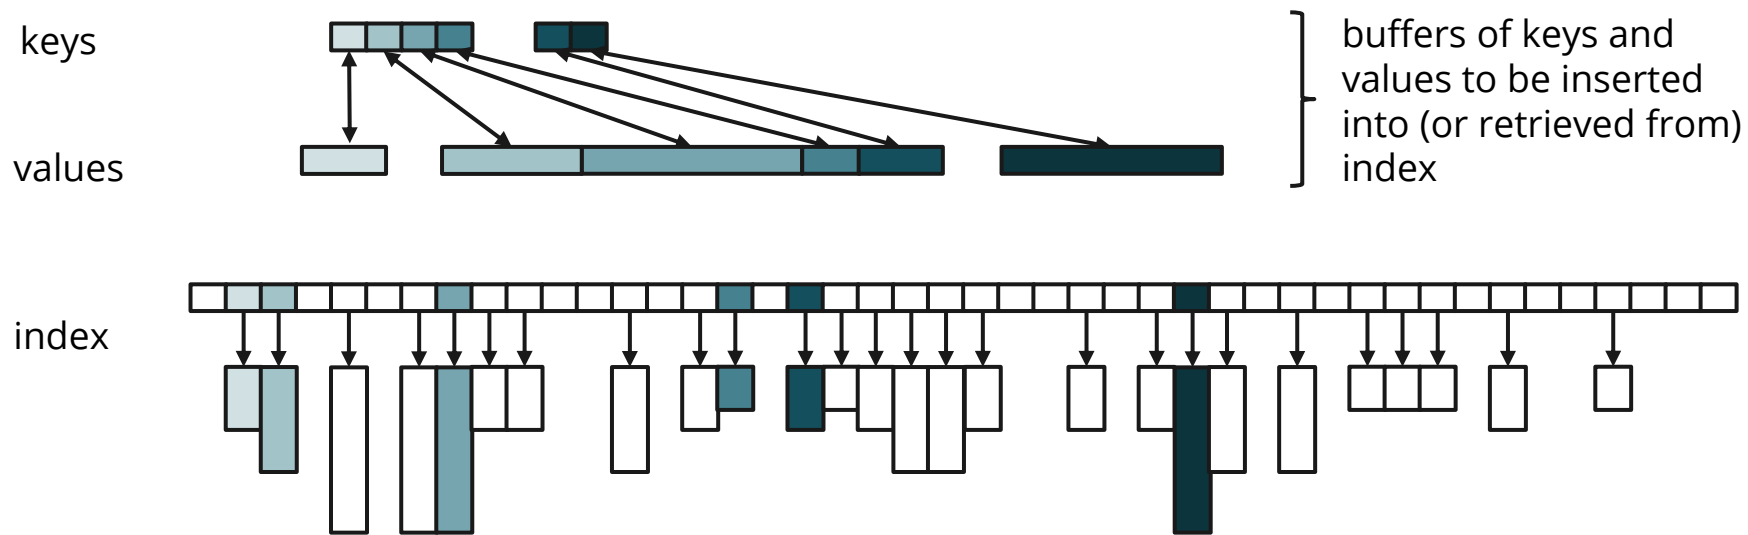
```
Motr_client_read(op, obj, data_buf_vec, attr_buf_vec, extent_vec)
```

```
Motr_client_write(op, obj, tx, data_buf_vec, attr_buf_vec, extent_vec)
```

```
Motr_client_alloc(op, obj, tx, extent_vec)
```

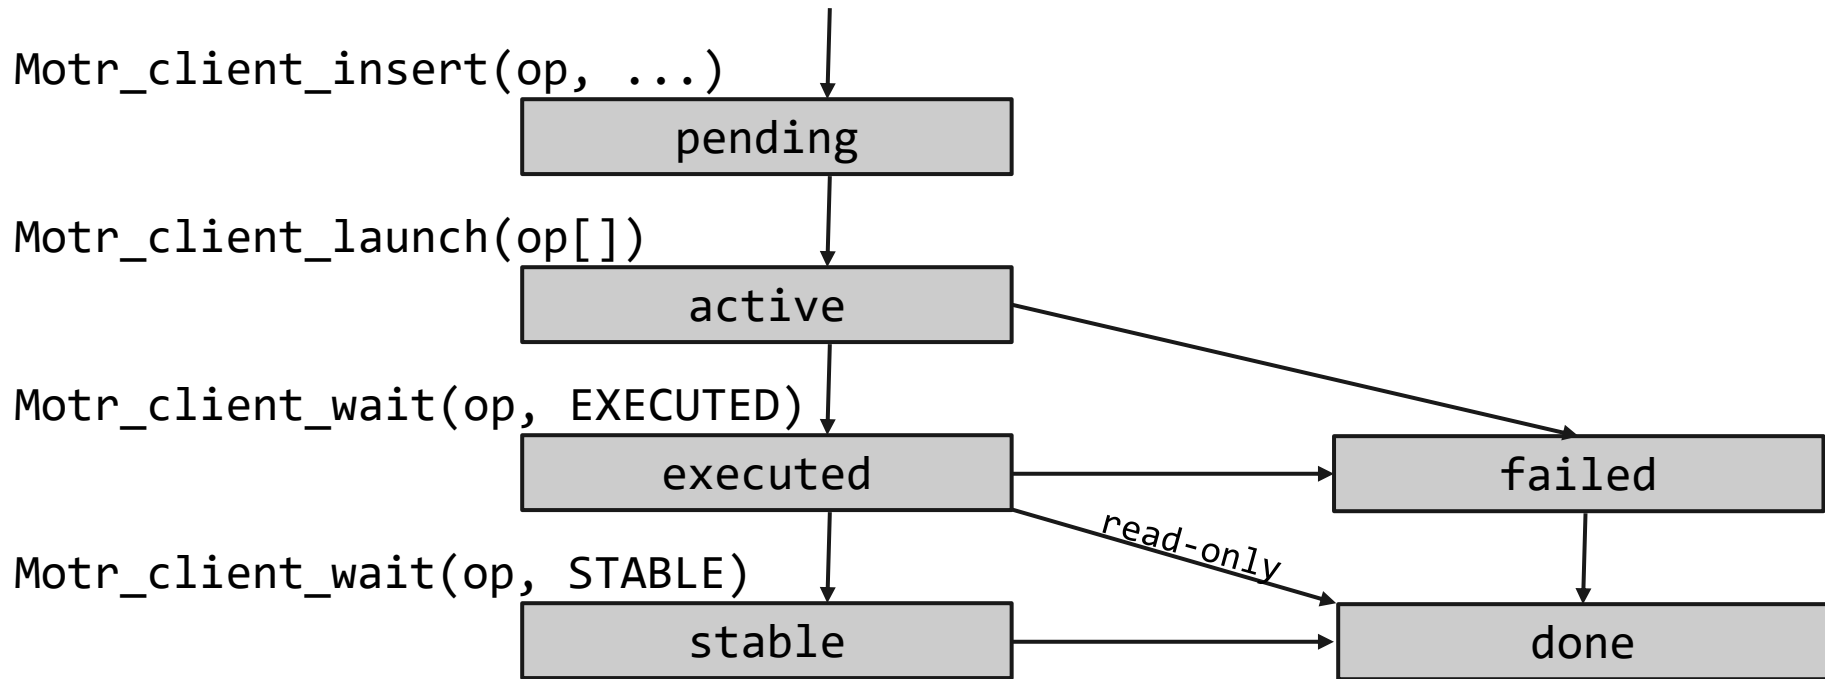
```
Motr_client_free(op, obj, tx, extent_vec)
```

Motr client: index



```
Motr_client_lookup(op, index, key_buf_vec, val_buf_vec)
Motr_client_insert(op, index, tx, key_buf_vec, val_buf_vec)
Motr_client_next(op, index, key_buf_vec, val_buf_vec)
```

Motr client: operation



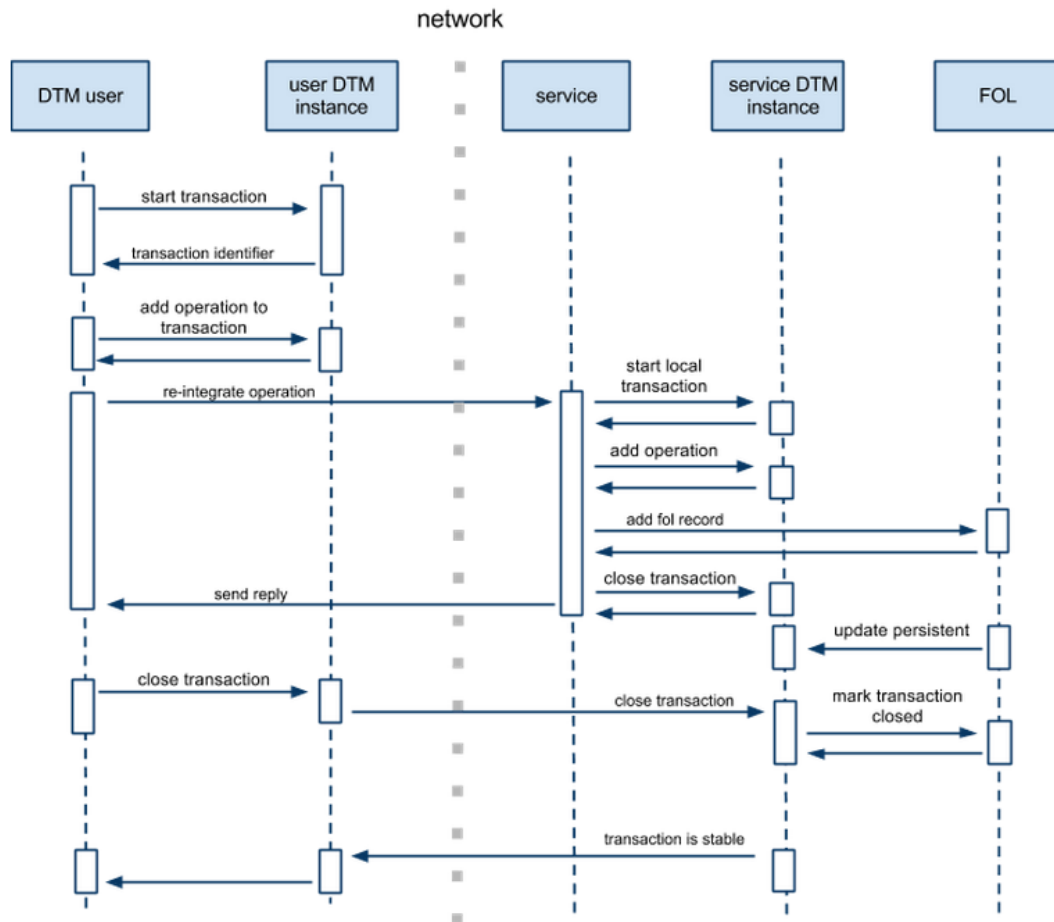
components

- Motr client
- **transactions (dtm)**
- resource manager
- fdmi
- addb
- network raid, layouts
- containers
- function shipping
- lingua franca
- integrity checking
- meta-data back-end
- network, rpc, fop, HA
- fom, reqh
- device io (stob)
- network raid repair
- security

dtm: transactions

- a transaction is a group of operations
- **a**tomicity w.r.t. certain failures (network partitions, node restarts)
- **c**onsistency is defined by user
- guarantees neither **i**solation nor serialisability
- doesn't guarantee synchronous **d**urability: too expensive for small transactions. Asynchronous stabilisation: the user is notified when the transaction becomes stable
- stabilisation ordering, liveness

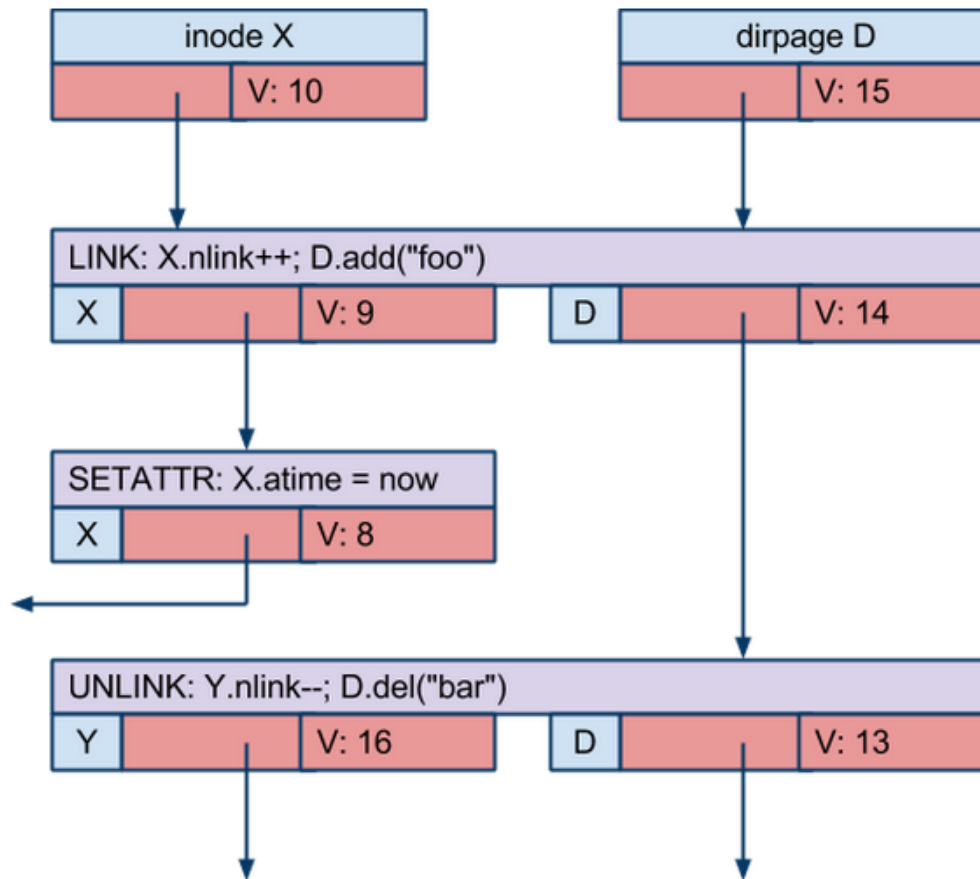
dtm: sequence



dtm: features

- masks certain transient failures
 - network partition, re-ordering, duplication
 - node restart
- write-ahead logging on each server
- undo for data
- redo for meta-data
- stabilisation: global logical clock (epochs)

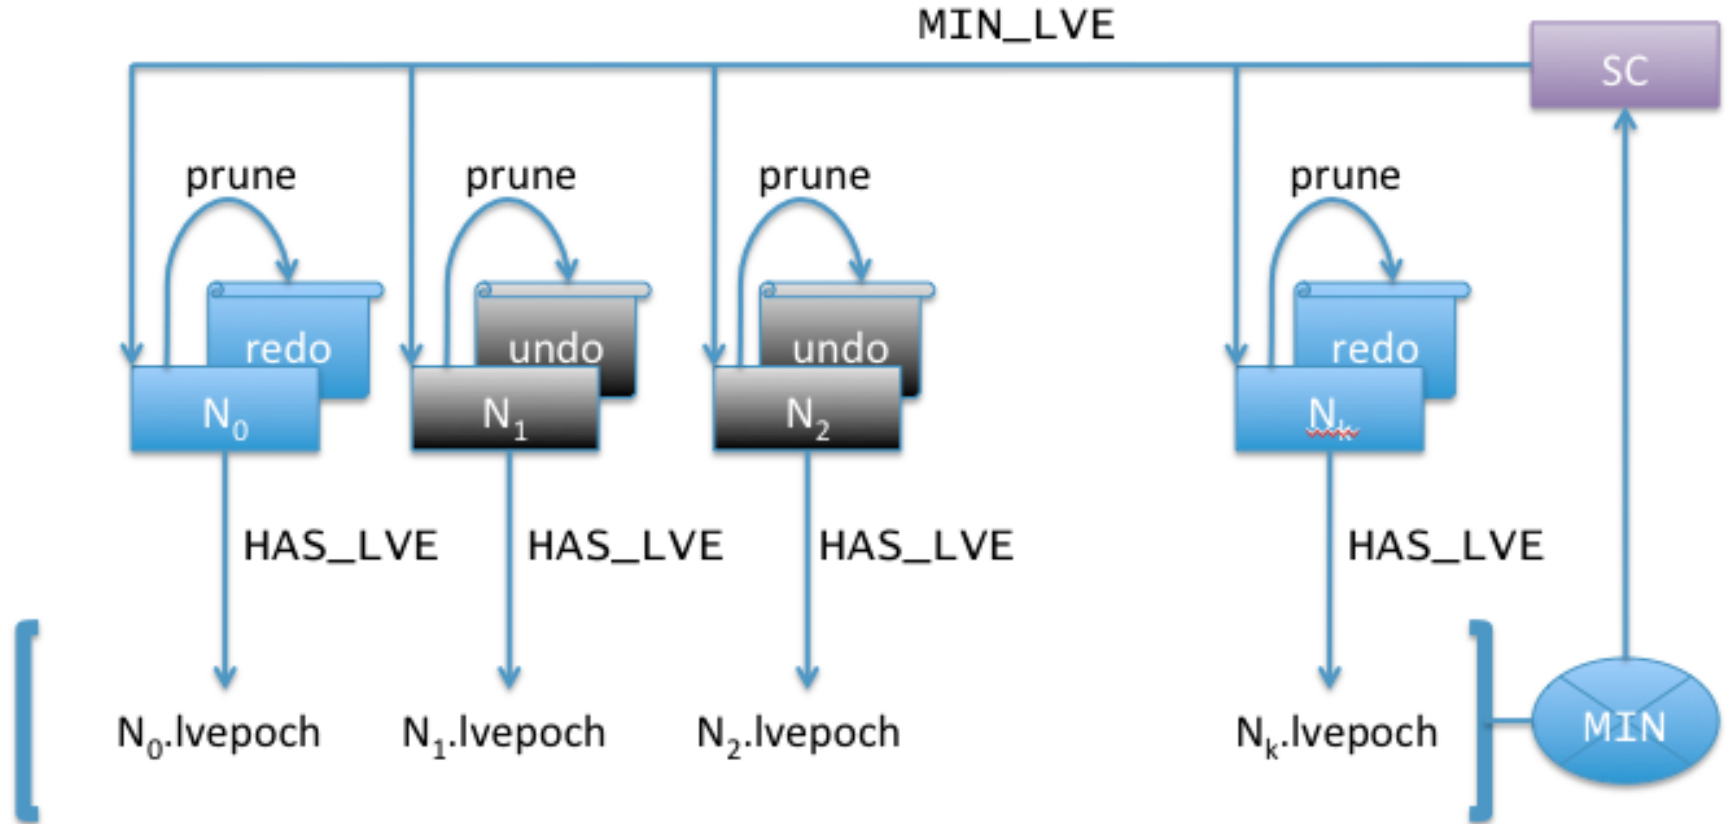
dtm: versions



dtm: epochs

- distributed clock to detect operation dependency and ordering (epoch, Fidge-Mattern, Lamport)
- messages are tagged with the logical timestamps:
Event1 depends on Event0, then $\text{Event0.epoch} \leq \text{Event1.epoch}$
- any node can advance its clock independently
- operations are kept in the persistent log in epoch order, until epoch is stable. Then the log is pruned
- global coordination to determine when an epoch is stable

dtm: epochs



components

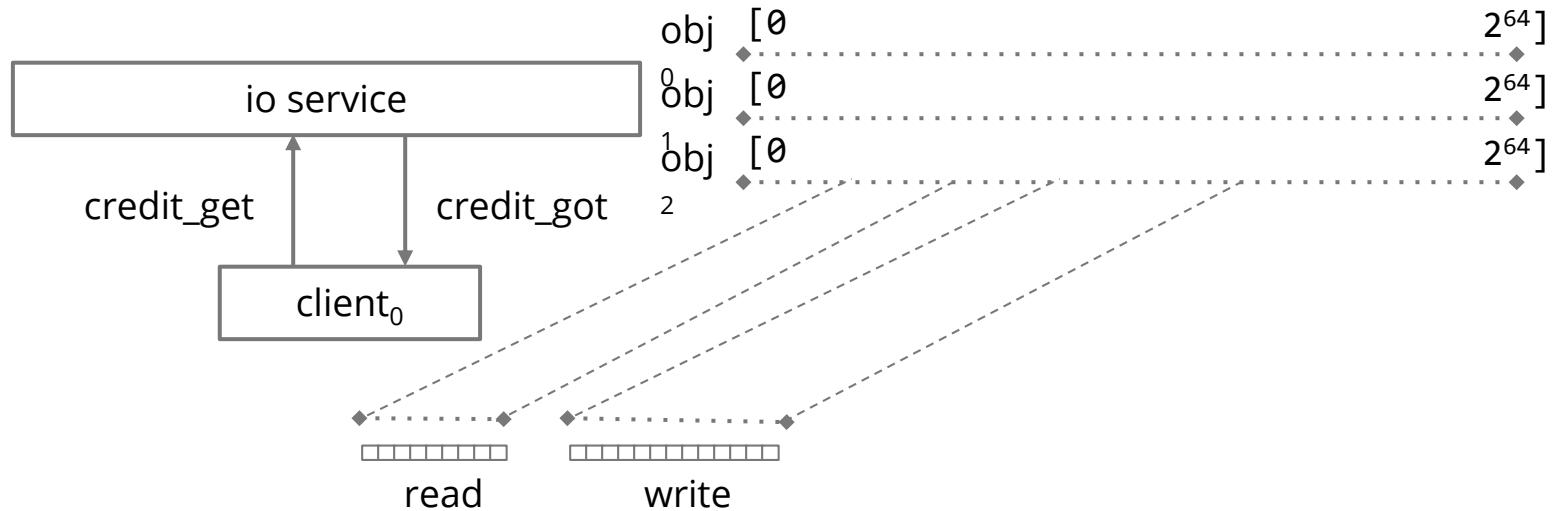
- Motr client
- transactions (dtm)
- **resource manager**
- fdmi
- addb
- network raid, layouts
- containers
- function shipping
- lingua franca
- integrity checking
- meta-data back-end
- network, rpc, fop, HA
- fom, reqh
- device io (stob)
- network raid repair
- security

rm: definition

- resource: anything with ownership. An extent in an object, an entire object, a key in an index, *etc.*
- credit: a right to use a resource in a particular way
- credits control:
 - distributed caching
 - concurrency
- credits can be borrowed and sublet
- resource manager is separate from resource
- resource manager resolves conflicts
- user can define new resource types

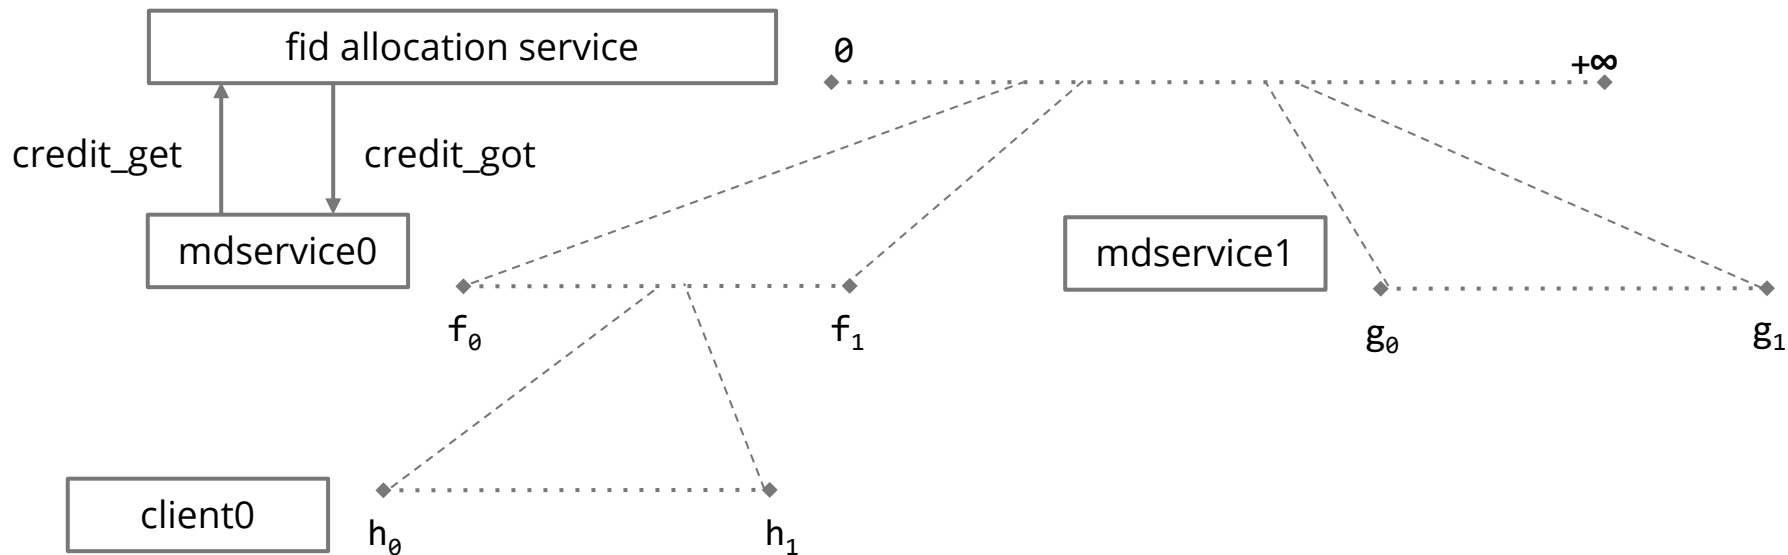
rm: use case

Example: Block extent in a Motr client object



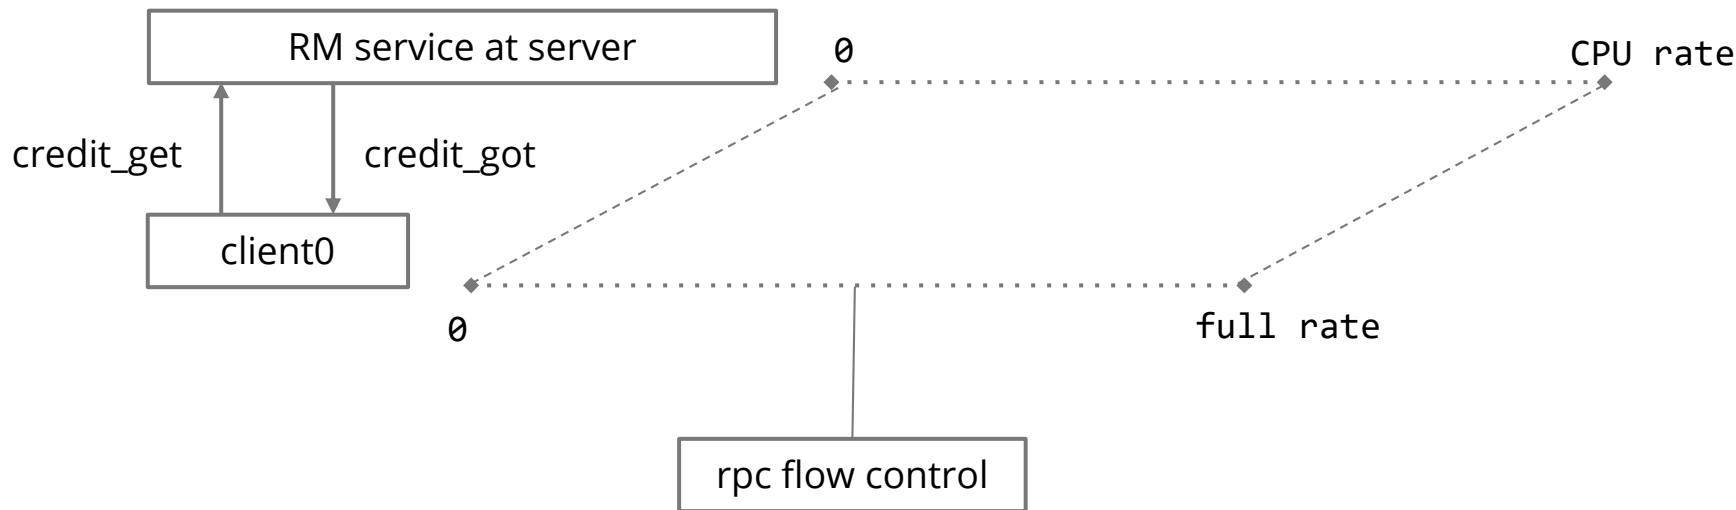
rm: use case

Example: fid extent allocation. Fid: 128 bit.



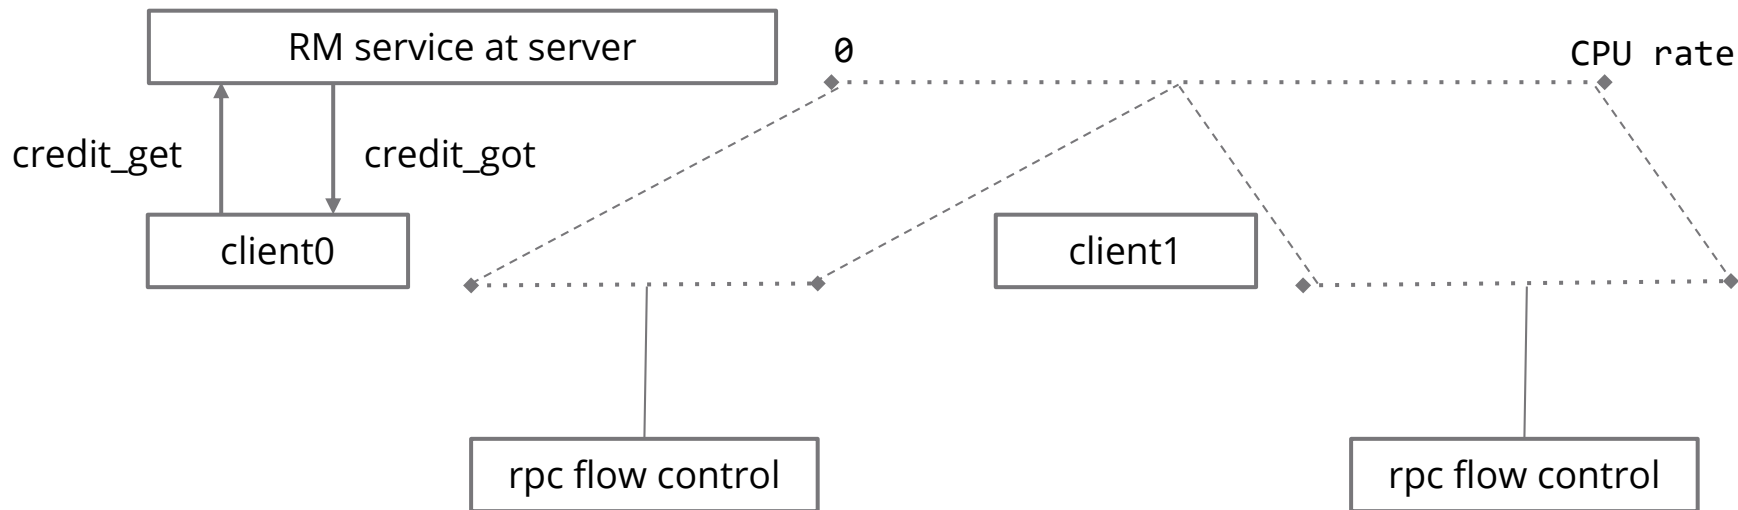
rm: use case

Example: server CPU cycles

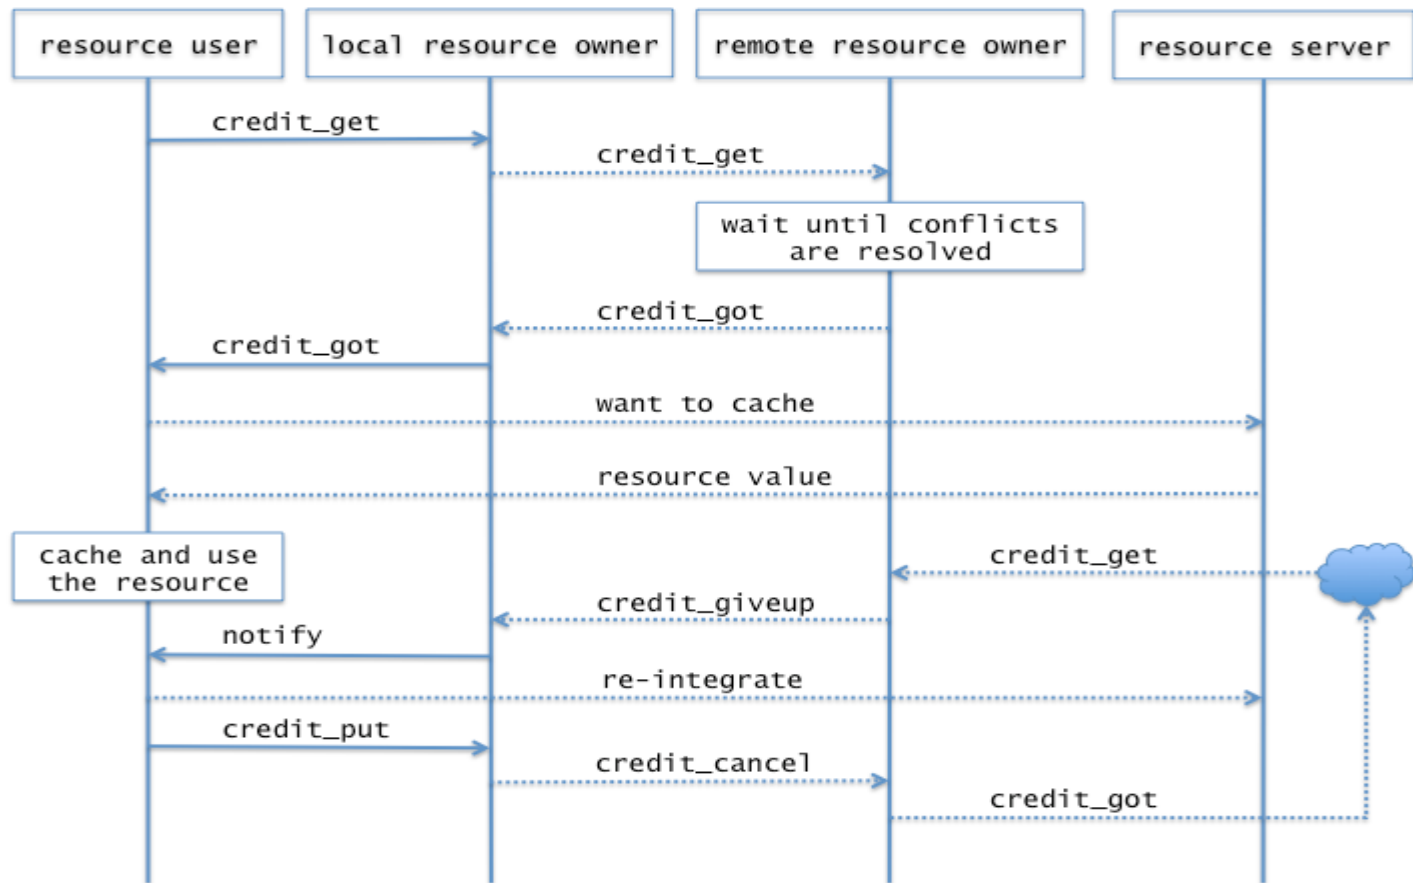


rm: use case

Example: server CPU cycles



rm: sequence



rm: resources

Resource types

- open-files
- file extents
- disk space (grants)
- quotas
- server memory
- server cpu cycles
- file identifiers (fids)
- inode numbers
- network bandwidth
- storage bandwidth
- layout
- cluster configuration
- power

rm: resource manager context

- generic infrastructure:
 - RM service: BORROW, REVOKE, CANCEL
 - client (Motr client) interface
- specific resource types:
 - resource and credit names
 - conflict, credit ordering
- RM users:
 - resource acquisition and release logic
 - cache invalidation
 - assignment of RM services

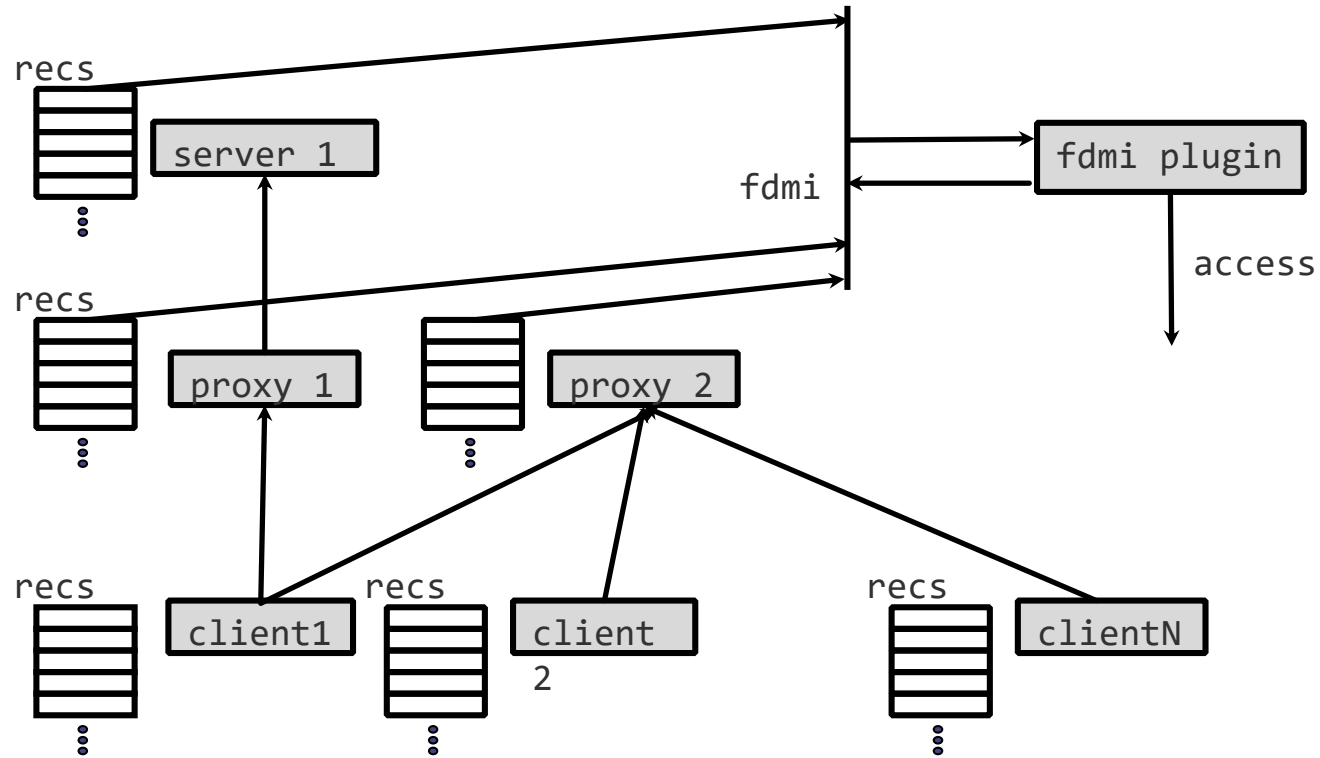
components

- Motr client
- transactions (dtm)
- resource manager
- **fdmi**
- addb
- network raid, layouts
- containers
- function shipping
- lingua franca
- integrity checking
- meta-data back-end
- network, rpc, fop, HA
- fom, reqh
- device io (stob)
- network raid repair
- security

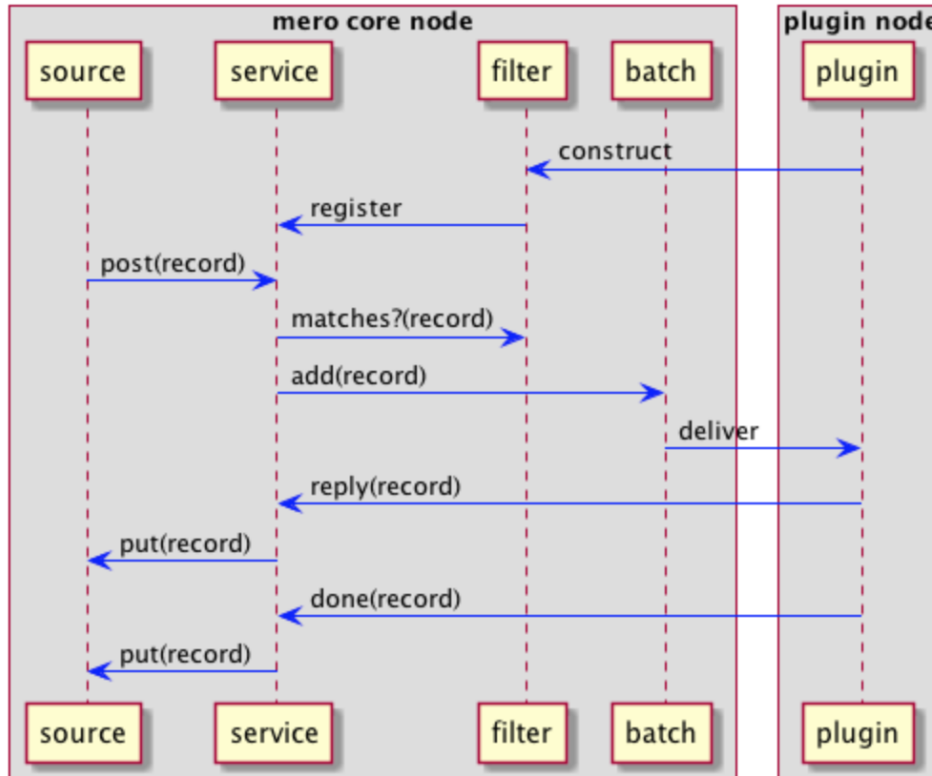
fdmi

- operation log (fol):
 - record each operation on an object
 - log consists of records, log maintained by each node
- file-system definition and manipulation interface (fdmi)
 - scalable publish-subscribe interface
 - subscribe to records matching certain filter
 - map-reduce-style mechanism to deliver matching records to the subscribers
 - transactional delivery (EOS)
 - delivers: fol records, addb records, HA events, others
- horizontal scalability
 - offload plugin processing and data-structures
 - asynchronous processing, batching

fdmi

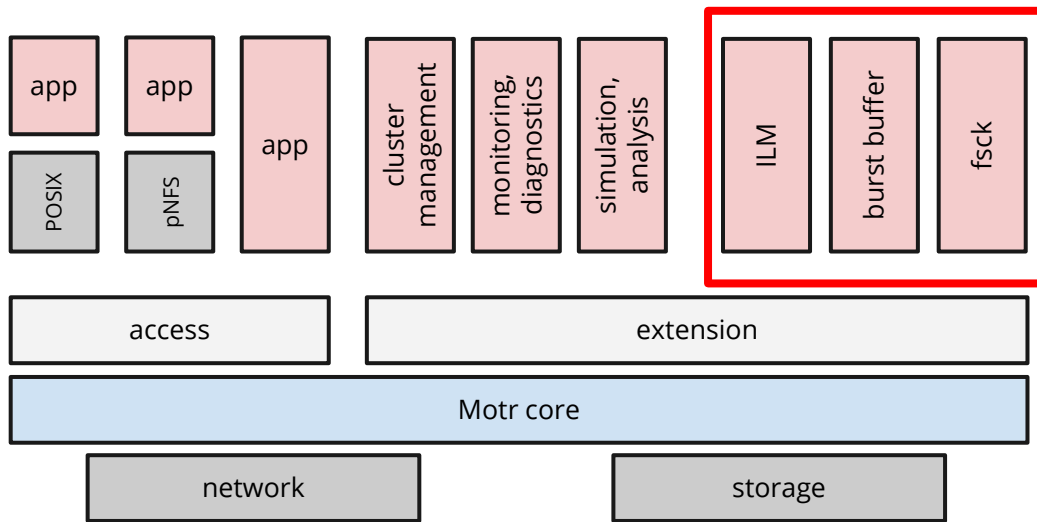


fdmi



fdmi

- ILM
 - replication
 - migration
 - backup, archival,
 - hsm
- indexing
- fsck
- data re-structuring
 - proxy de-staging
 - RAID re-striping
- guided interfaces
 - profiling
 - prefetching, destaging



components

- Motr client
- transactions (dtm)
- resource manager
- fdmi
- **adddb**
- network raid, layouts
- containers
- function shipping
- lingua franca
- integrity checking
- meta-data back-end
- network, rpc, fop, HA
- fom, reqh
- device io (stob)
- network raid repair
- security

addb

- systems grow larger and more complex
- how well the system is utilised?
- is it failure or expected behaviour?
- is it system or application behaviour?
- sources of data:
 - system logs
 - operating system
 - application traces
- very large amount of collected data
- ... or insufficiently detailed, or both
- difficult to analyse and correlate

addb

- instrumentation on client and server
- data about operation execution and system state
- passed through network
- cross-referenced
- always on (post-mortem analysis, first incident fix)
- simulation (change configuration, larger system, load mix)

```
* 2015-04-20-14:36:13.687531192 alloc size: 40, addr: @0x7fd27c53eb20
| node          <f3b62b87d9e642b2:96a4e0520cc5477b>
| locality      1
| thread        7fd28f5fe700
| fom           @0x7fd1f804f710, 'IO fom' transitions: 13 phase: Zero-copy finish
| stob-io-launch 2015-04-20-14:36:13.629431319, <200000000000003:10000>, count: 8, bvec-nr: 8, ivec-nr: 1, offset: 0
| stob-io-launch 2015-04-20-14:36:13.666152841, <100000000adf11e:3>, count: 8, bvec-nr: 8, ivec-nr: 8, offset: 65536
```

addb: anatomy of a record

```
* 2015-04-14-15:33:11.998165453 fom-descr service: <7300000000000001:0>, sender: c28baccf27e0001, req-opcode: Read request,
  rep-opcode: Read reply, local: false
  |   node           <11186d8bf0e34117:ab1897c062a22573>
  |   locality       3
  |   thread         7f79e57fb700
  |   ast
  |   fom            @0x7f795008ed20, 'IO fom', transitions: 0, phase: 0
```

- measurement and context
- timestamped
- labels: identify context
- payload: up to 16 64-bit values,
- interpreted by consumer

addb: sensors and histograms

```
* 2018-04-05-14:27:40.563070315 fom-active nr: 710 min: 0 max: 24 avg: 5.029577 dev: 11.147012
1 1: 226 3: 313 5: 98 7: 26 9: 20 11: 8 13: 5 15: 6 17: 4 19: 3 21: 0 23: 0 25: 0
|
| node <5e341757d0cf46eb:92a6d6e991b46387>
| pid 6627
| locality 0
```

- some events are too frequent
- collapse them into counters
- count last events per locality
- automatically size buckets

```
      : 1 |
1 : 226 | *****
3 : 313 | *****
5 : 98  | *****
7 : 26  | ***
9 : 20  | **
11 : 8   | *
13 : 5   |
15 : 6   |
17 : 4   |
19 : 3   |
21 : 0   |
23 : 0   |
25 : 0   |
```

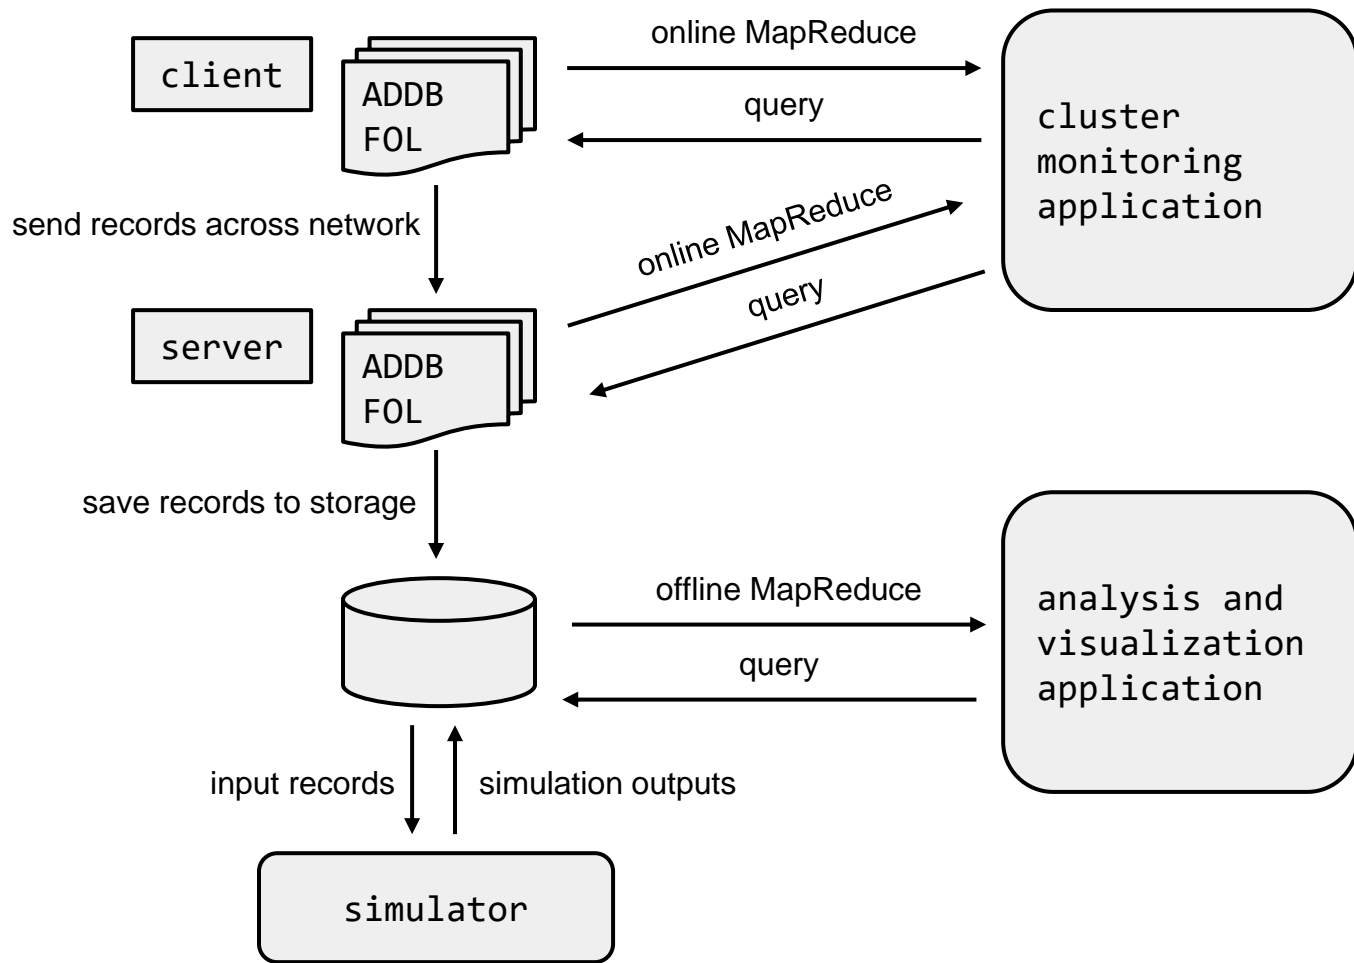
addb: state-machine transitions

```
* 2018-04-05-14:37:09.660378140 be-op/m0_be_op::bo_sm: M0_BOS_ACTIVE -[completed]-> M0_BOS_DONE
nr: 27788 min: 0 max: 556380 avg: 785.763279 dev: 56005939.461453 datum: 0
0 0: 25862 2066: 1138 4132: 309 6198: 98 8264: 61 10330: 67 12396: 68
  14462: 30 16528: 12 18594: 10 20660: 23 22726: 29 24792: 81
|
| node <5e341757d0cf46eb:92a6d6e991b46387>
| pid 6627
| locality 2
```

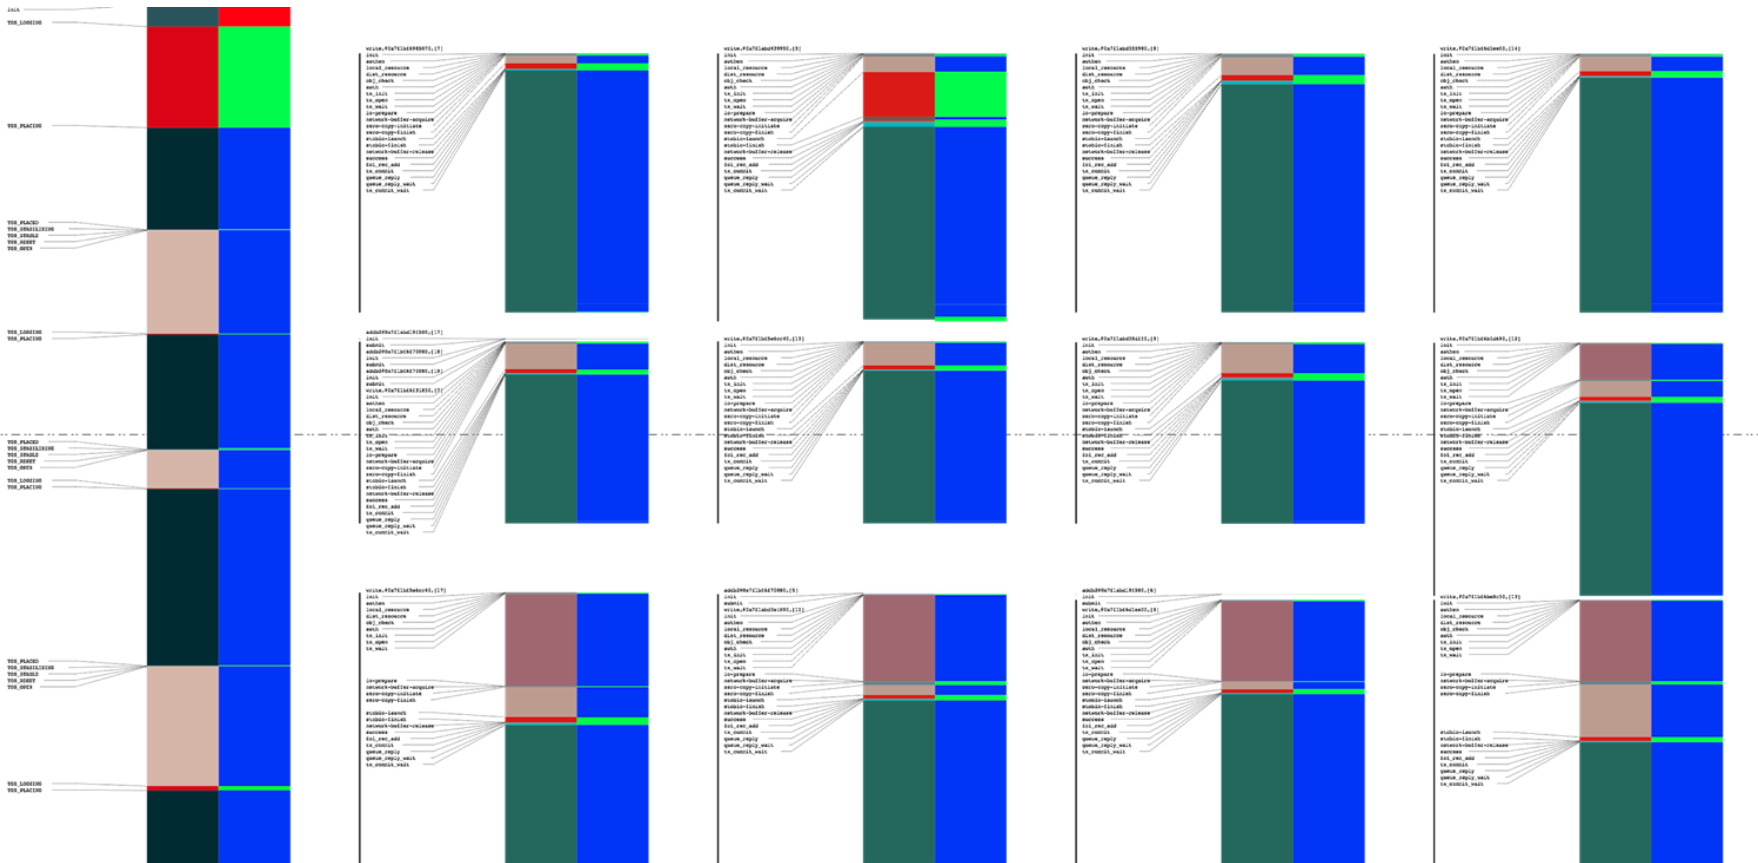
- request processing on clients and servers
- state machines
- state transition delays: in "binary microseconds", second >> 10

:	0	
0 :	25862	*****
2066 :	1138	*****
4132 :	309	***
6198 :	98	*
8264 :	61	*
10330 :	67	*
12396 :	68	*
14462 :	30	
16528 :	12	
18594 :	10	
20660 :	23	
22726 :	29	
24792 :	81	*

addb



addb

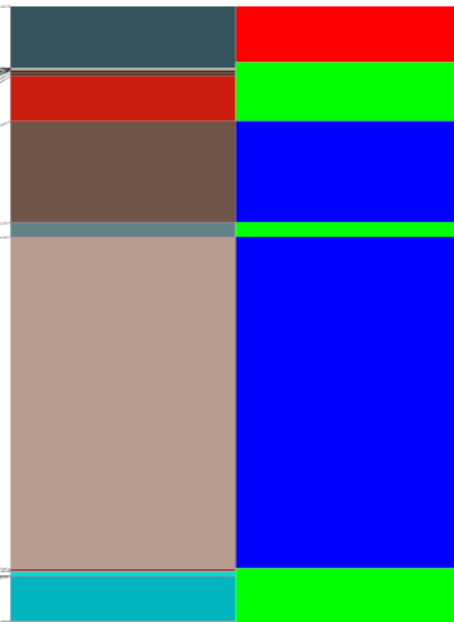


adddb

```
read,0x7f42e81e6df0,
init
```

```
authen
local_resource
dist_resource
obj_check
auth
tx_init
tx_open
tx_wait
io-prepare
network-buffer-acquire
stobio-launch
stobio-finish
zero-copy-initiate
zero-copy-finish
```

```
network-buffer-release
success
fol_rec_add
tx_commit
queue_reply
queue_reply_wait
tx_commit_wait
```



```
L <200000100000016:17>
L <100000000000000:1>
E <100000000000000:1>
E <200000100000016:17>
```

```
L <20000000000000016:17>
L <200000200000016:17>
L <200000300000016:17>
L <100000000000000:2>
E <100000000000000:2>
E <200000200000016:17>
L <100000000000000:3>
L <100000000000000:0>
L <200000400000016:17>
L <100000000000000:4>
E <100000000000000:3>
E <200000300000016:17>
E <100000000000000:0>
E <200000000000000:1>
E <100000000000000:4>
E <200000400000016:17>
```

```
Q #0x7f13a011f440 msg-send 200
```

```
Q #0x7f13903cf4b0 msg-send 200
```

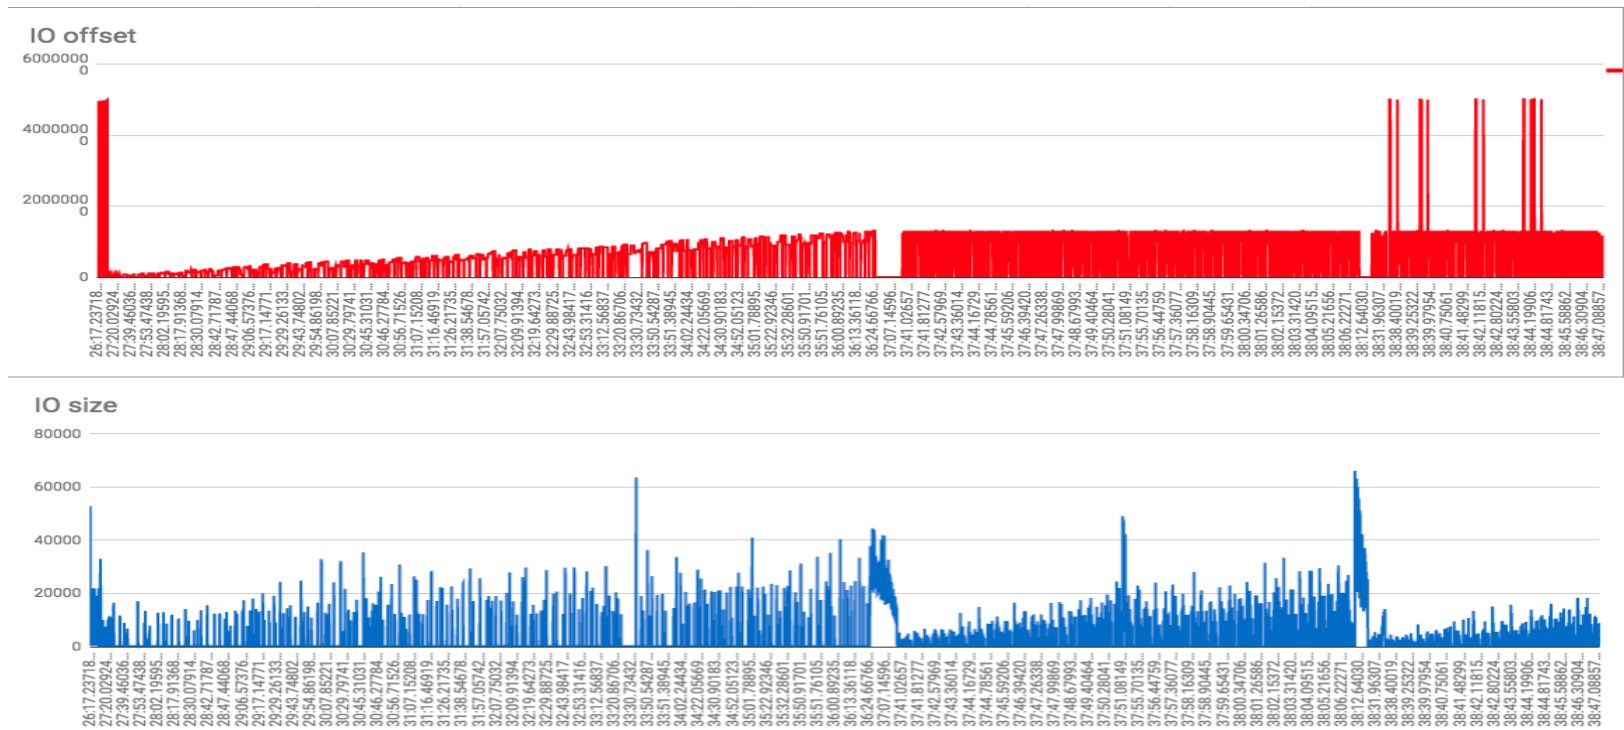
```
C #0x7f13a011f440
C #0x7f13903cf4b0
```

```
Q #0x7f1398052d70 msg-send 200
Q #0x7f13b005bbf0 a-bulk-recv 524288
Q #0x7f13b0055d70 a-bulk-recv 524288
Q #0x7f13b005b820 a-bulk-recv 524288
Q #0x7f13b0055c00 a-bulk-recv 524288
Q #0x7f13b0046f70 a-bulk-recv 524288
Q #0x7f13b005ba60 a-bulk-recv 524288
Q #0x7f13b0046de0 a-bulk-recv 1048576
```

```
C #0x7f1398052d70
C #0x7f13b005bbf0
C #0x7f13b0055d70
C #0x7f13b005b820
C #0x7f13b0055c00
C #0x7f13b0046f70
C #0x7f13b005ba60
C #0x7f13b0046de0
```

adadb: ad hoc profiling

```
$ m0adadb2dump ... | grep 'stob-io-launch' | awk '{print $2, $6, $12}'
```



addb: interface

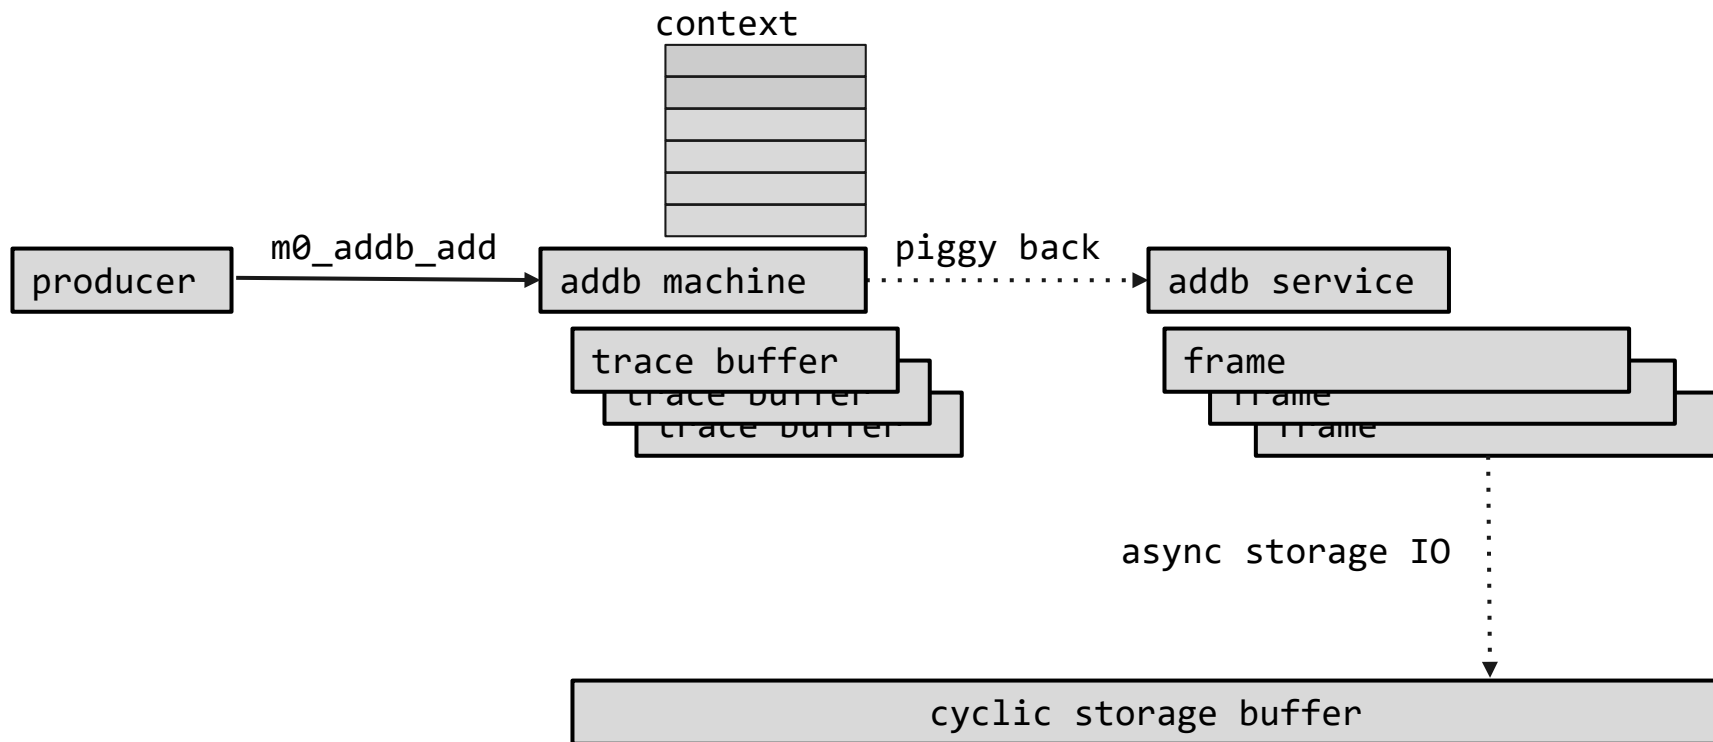
```
/**
 * Adds a label to the current context.
 *
 * @param id    - label identifier
 * @param n     - number of 64-bit values in label payload
 * @param value - payload
 */
void m0_addb_push(uint64_t id, int n, const uint64_t *value);

/**
 * Removes the top-most label in the current context stack.
 *
 * @param id - label identifier
 *
 * @pre "id" must be the identifier of the top-most label.
 */
void m0_addb_pop(uint64_t id);

/**
 * Adds one-time measurement in the current context.
 *
 * @param id    - measurement identifier
 * @param n     - number of 64-bit values in measurement payload
 * @param value - payload
 */
void m0_addb_add(uint64_t id, int n, const uint64_t *value);
```

- very simple interface
- binary values only
- stack (LIFO) context
- context management is modular
- separate interface for sensors (not shown)
- usable from system and applications

addb: data flow



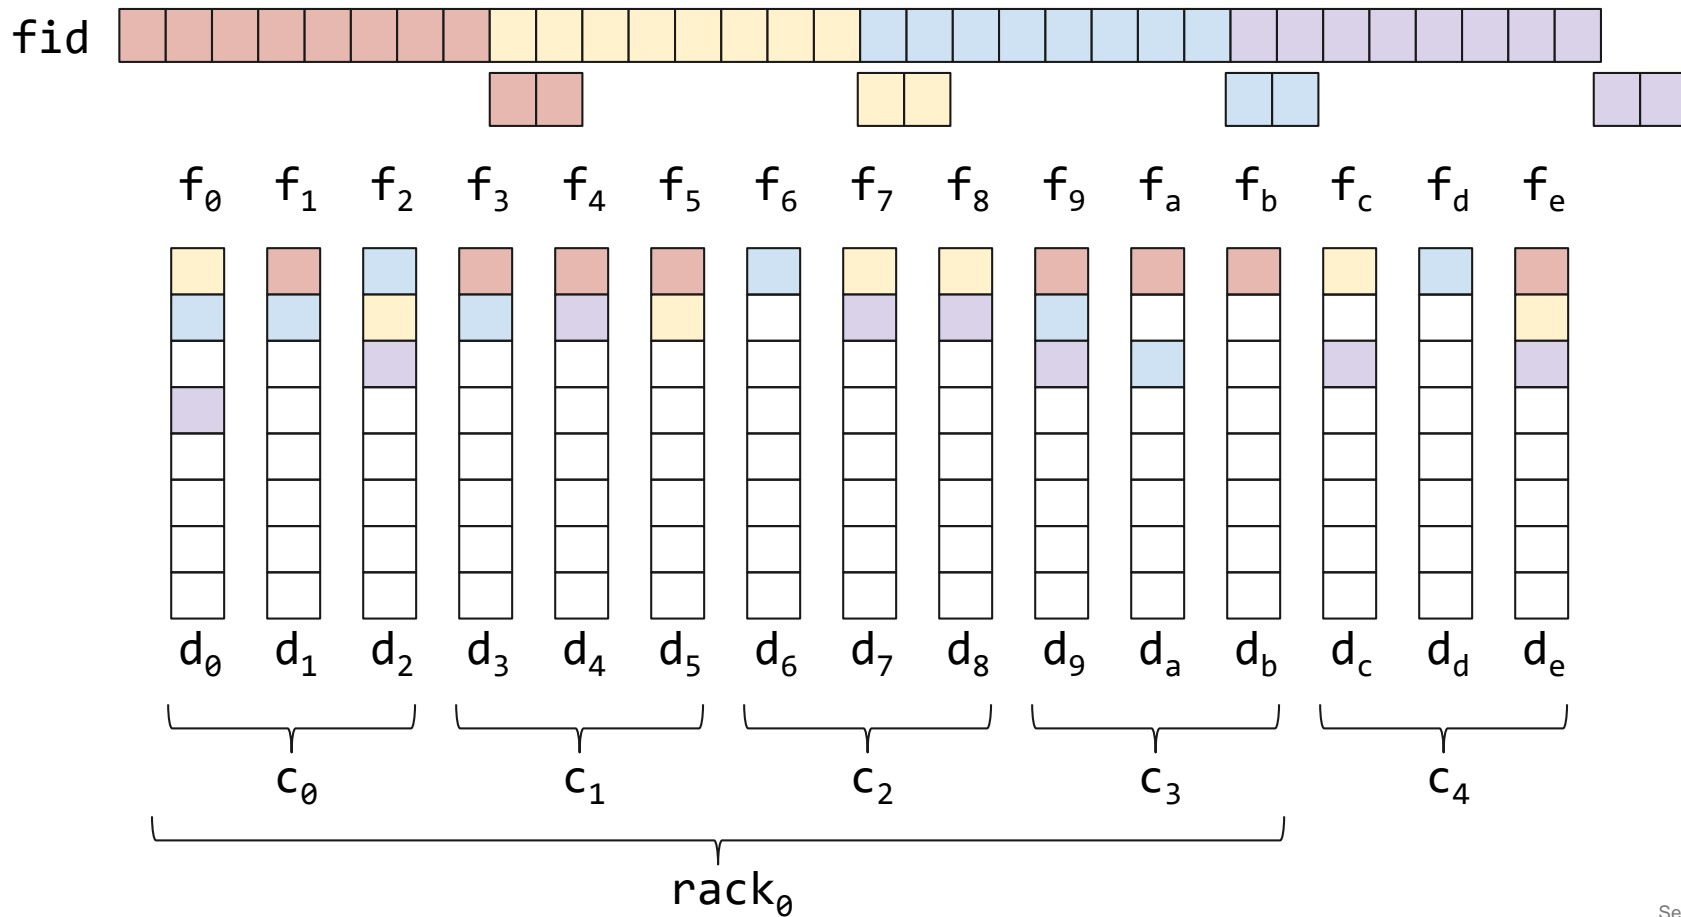
components

- Motr client
- transactions (dtm)
- resource manager
- fdmi
- addb
- **network raid, layouts**
- containers
- function shipping
- lingua franca
- integrity checking
- meta-data back-end
- network, rpc, fop, HA
- fom, reqh
- device io (stob)
- network raid repair
- security

io: layout

- determines how an object is stored in underlying containers
- layouts for data and meta-data
- examples:
 - network striping with parity de-clustering (default)
 - compression
 - encryption
 - de-duplication
 - composite (NBA, small files, migration)

io: parity groups



io: permutations

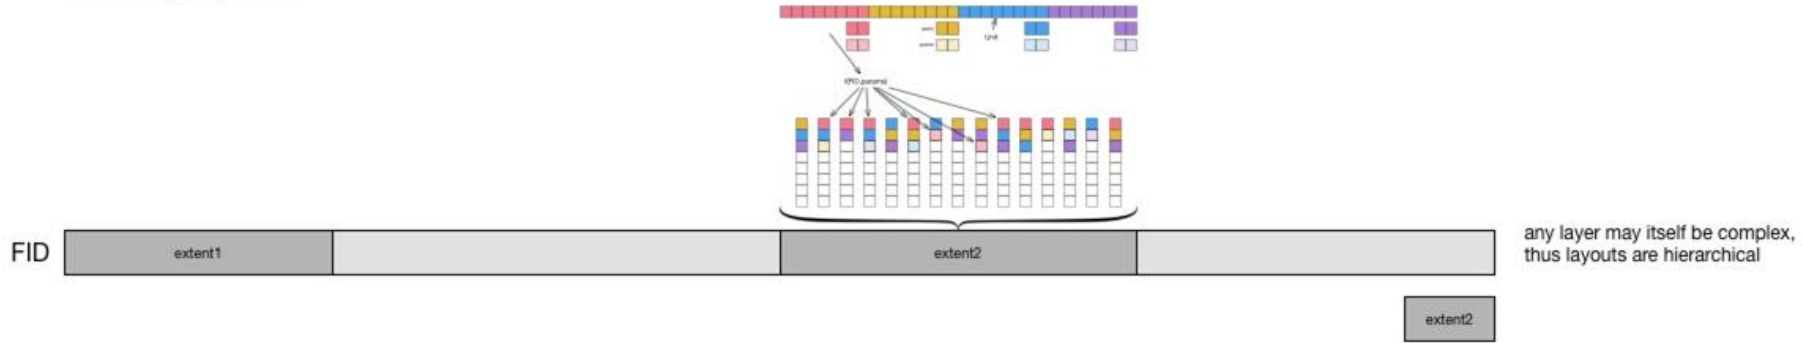
0 ₀	0 ₁	0 ₂	0 ₃	0 ₄	0 _p	0 _q	1 ₀	1 ₁	1 ₂	1 ₃	1 ₄	1 _p	1 _q	2 ₀
2 ₁	2 ₂	2 ₃	2 ₄	2 _p	2 _q	3 ₀	3 ₁	3 ₂	3 ₃	3 ₄	3 _p	3 _q	4 ₀	4 ₁
4 ₂	4 ₃	4 ₄	4 _p	4 _q	5 ₀	5 ₁	5 ₂	5 ₃	5 ₄	5 _p	5 _q	6 ₀	6 ₁	6 ₂
6 ₃	6 ₄	6 _p	6 _q	7 ₀	7 ₁	7 ₂	7 ₃	7 ₄	7 _p	7 _q	8 ₀	8 ₁	8 ₂	8 ₃
8 ₄	8 _p	8 _q	9 ₀	9 ₁	9 ₂	9 ₃	9 ₄	9 _p	9 _q	10 ₀	10 ₁	10 ₂	10 ₃	10 ₄
10 _p	10 _q	11 ₀	11 ₁	11 ₂	11 ₃	11 ₄	11 _p	11 _q	12 ₀	12 ₁	12 ₂	12 ₃	12 ₄	12 _p
12 _q	13 ₀	13 ₁	13 ₂	13 ₃	13 ₄	13 _p	13 _q	14 ₀	14 ₁	14 ₂	14 ₃	14 ₄	14 _p	14 _q
PDRAID [15 (5+2+0)], 1 Tile														

0 ₃	1 ₁	0 ₁	1 _q	1 ₂	0 _q	2 ₀	1 ₄	0 ₀	1 ₀	0 _p	1 _p	0 ₄	0 ₂	1 ₃
2 ₄	3 ₂	2 ₂	4 ₀	3 ₃	3 ₀	4 ₁	3 _p	2 ₁	3 ₁	2 _q	3 _q	2 _p	2 ₃	3 ₄
4 _p	5 ₃	4 ₃	6 ₁	5 ₄	5 ₁	6 ₂	5 _q	4 ₂	5 ₂	5 ₀	6 ₀	4 _q	4 ₄	5 _p
6 _q	7 ₄	6 ₄	8 ₂	7 _p	7 ₂	8 ₃	8 ₀	6 ₃	7 ₃	7 ₁	8 ₁	7 ₀	6 _p	7 _q
9 ₀	9 _p	8 _p	10 ₃	9 _q	9 ₃	10 ₄	10 ₁	8 ₄	9 ₄	9 ₂	10 ₂	9 ₁	8 _q	10 ₀
11 ₁	11 _q	10 _q	12 ₄	12 ₀	11 ₄	12 _p	12 ₂	10 _p	11 _p	11 ₃	12 ₃	11 ₂	11 ₀	12 ₁
13 ₂	14 ₀	13 ₀	14 _p	14 ₁	13 _p	14 _q	14 ₃	12 _q	13 _q	13 ₄	14 ₄	13 ₃	13 ₁	14 ₂
PDRAID [15 (5+2+0)], 1 Tile														

layout: composite

Complex Layout

individual layout per extent



layout: composite

Clovis Object

a single top-level "Object" in Mero

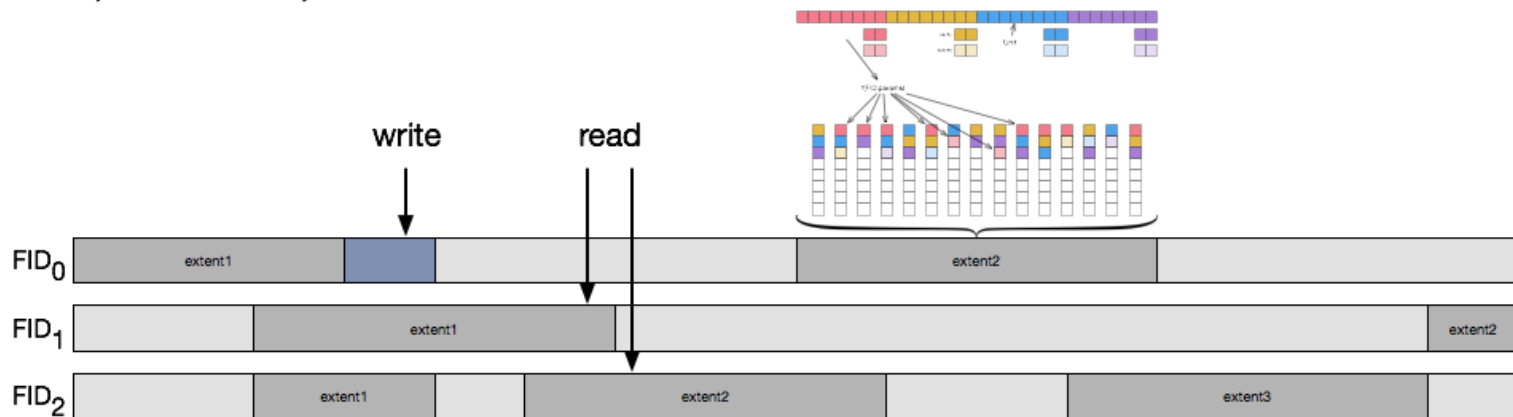
FID



Clovis block (IO) size is LCM of unit sizes in layout

Composite Layout

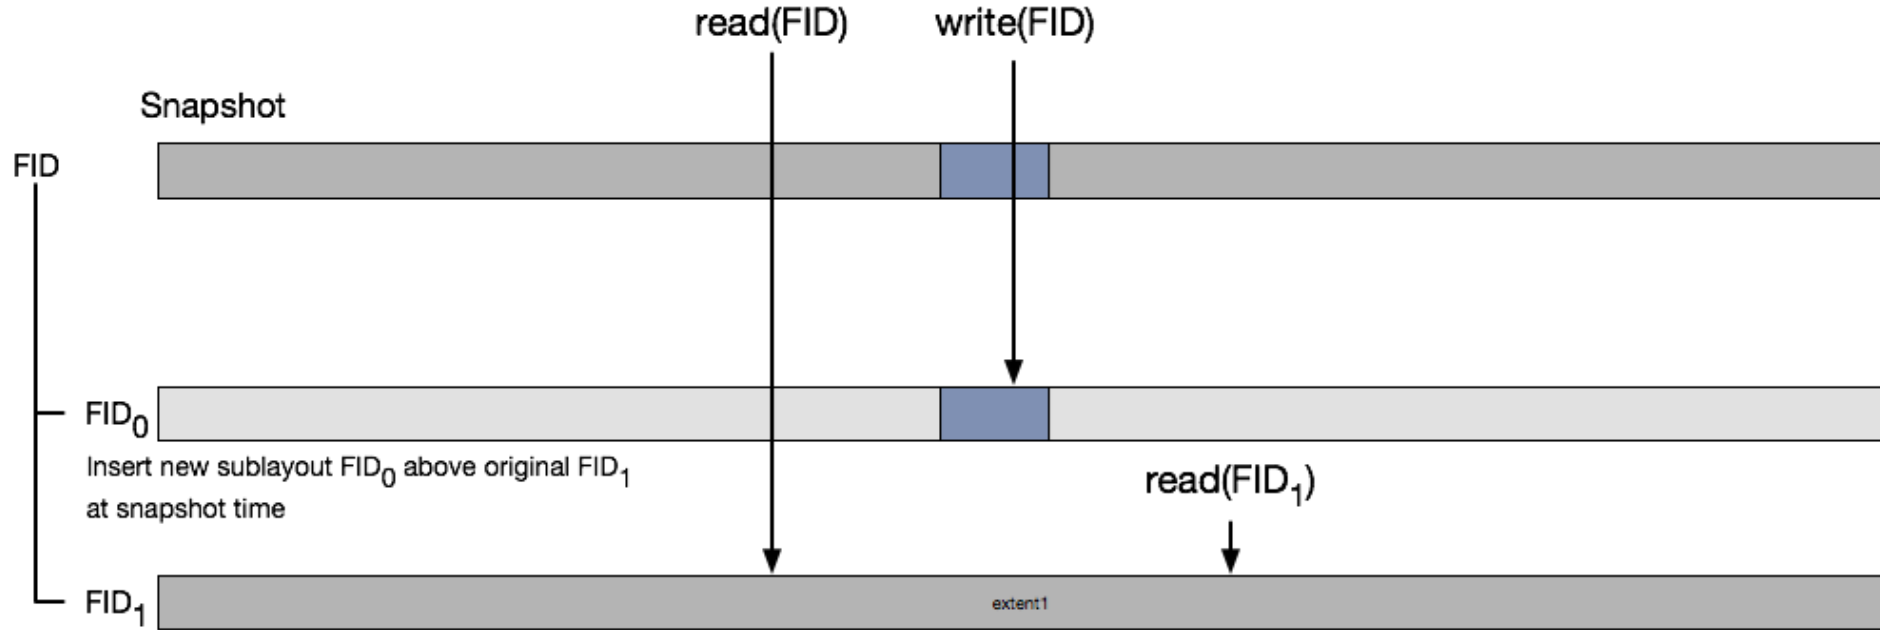
layered set of individual layouts



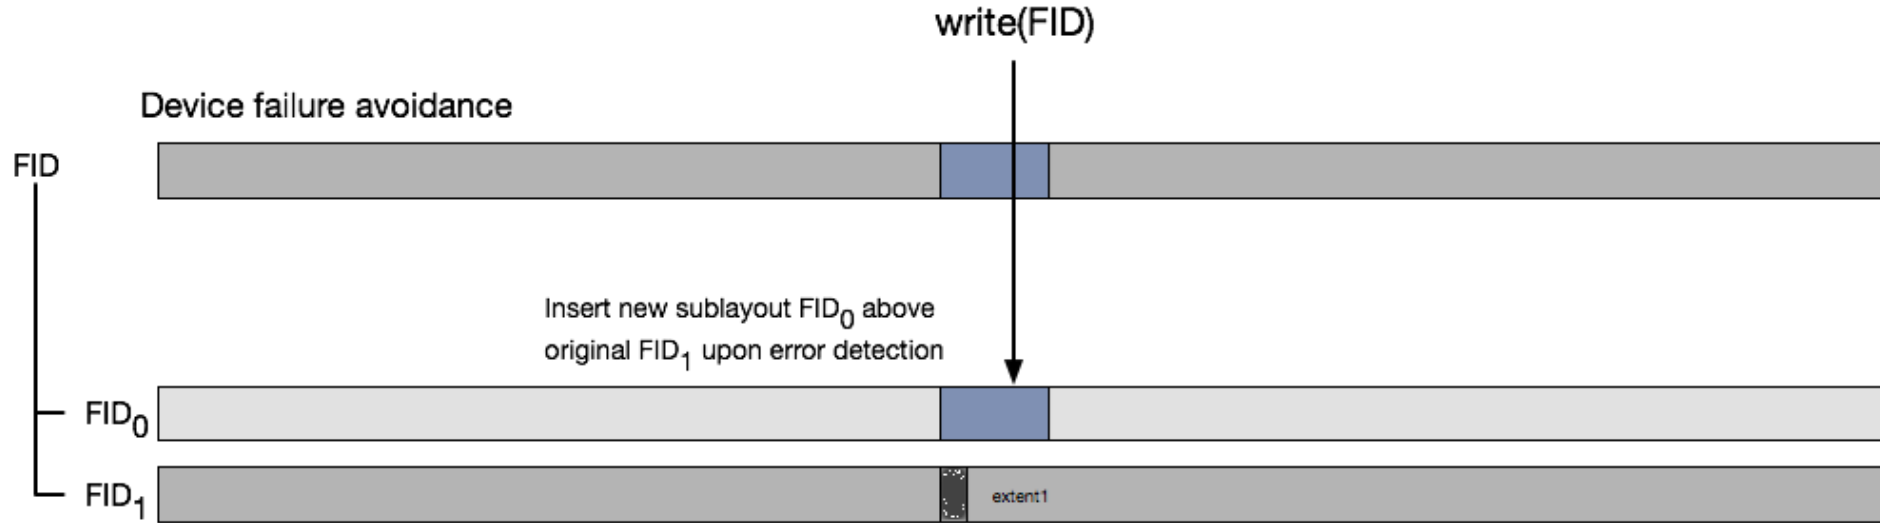
any layer may itself be complex,
thus layouts are hierarchical

Reads and writes fall through to first layer with mapped extents.
(Newly written blocks are added to read extents.) Layers can also be
read/written directly using FID_{sub}

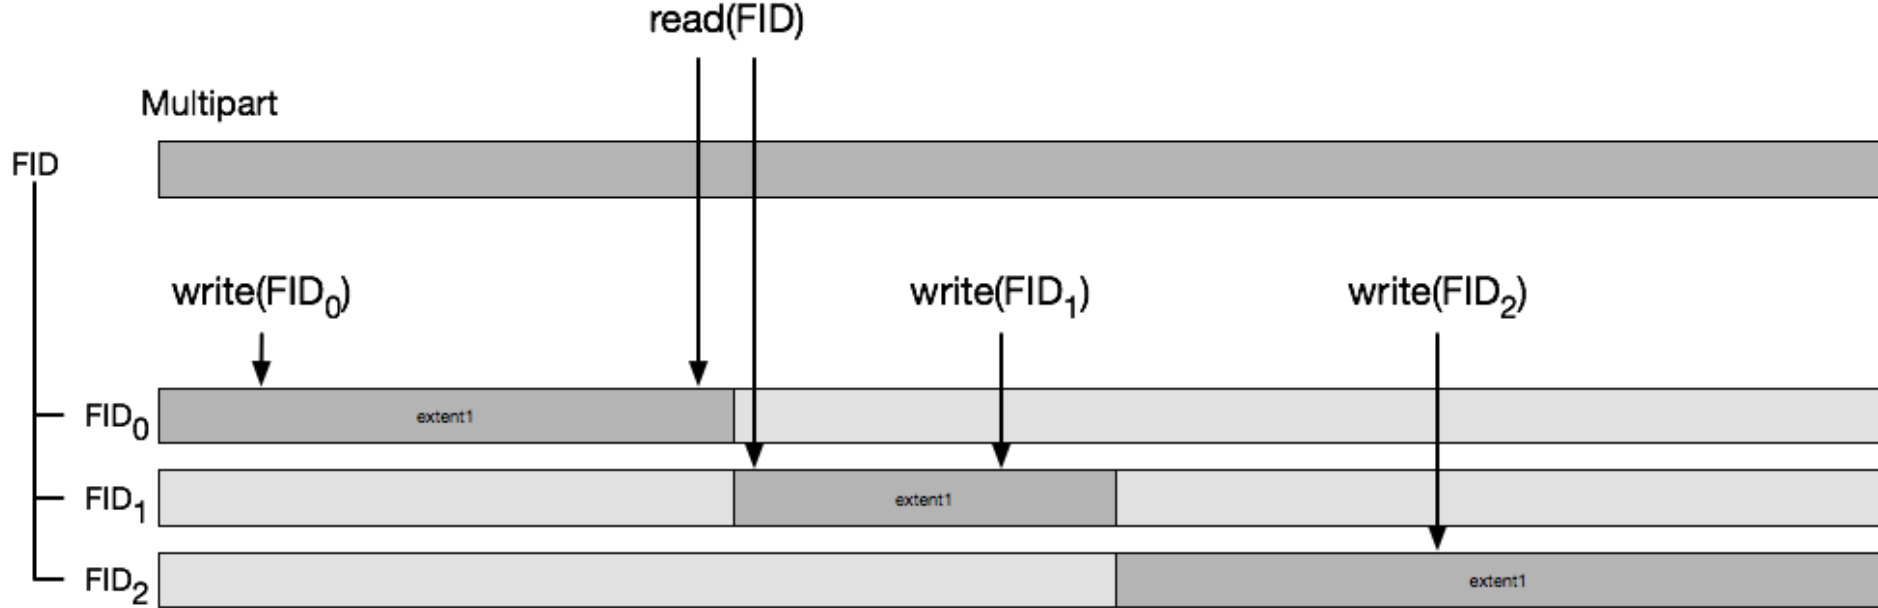
layout: composite: snapshot



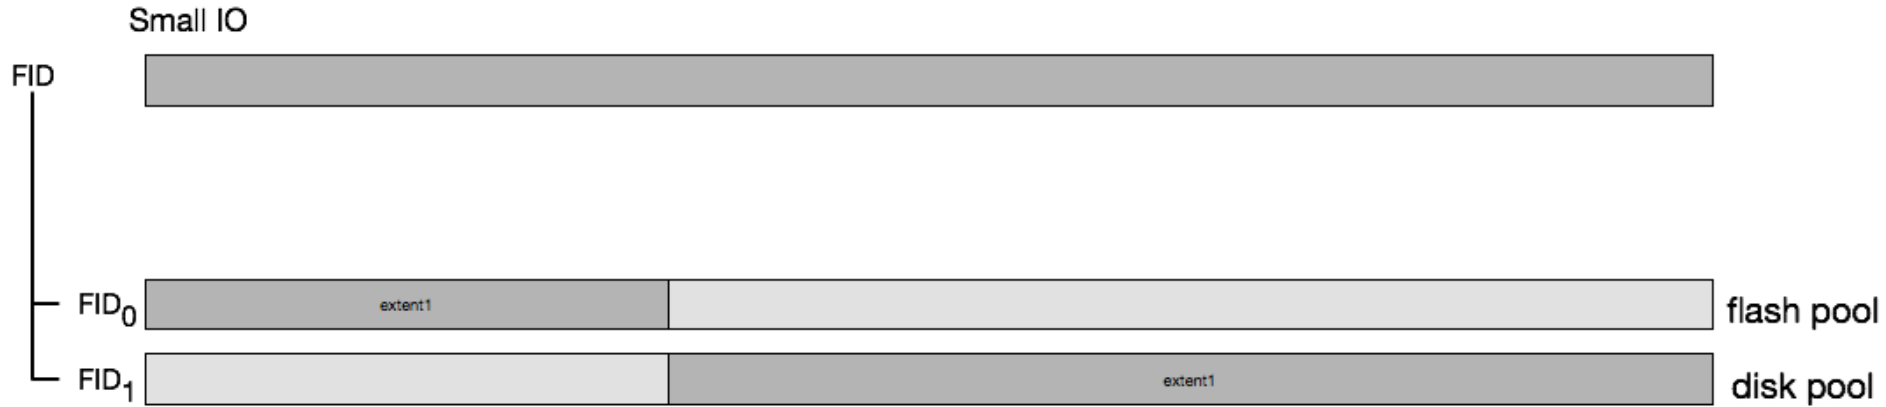
layout: composite: nba



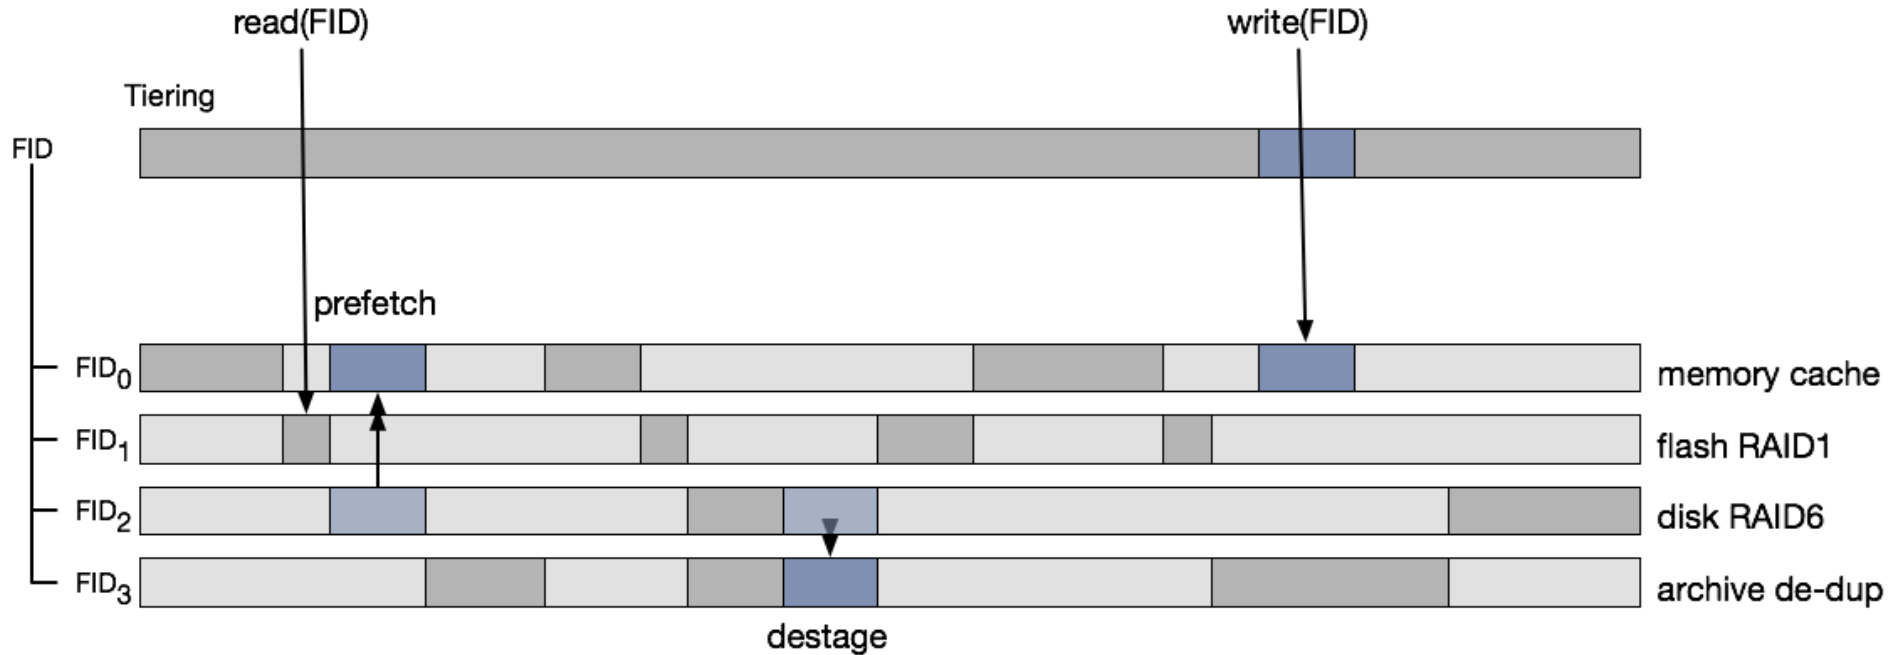
layout: composite: s3 partial upload



layout: composite: small files IO



layout: composite: HSM, tiering



components

- Motr client
- transactions (dtm)
- resource manager
- fdmi
- addb
- network raid, layouts
- **containers**
- function shipping
- lingua franca
- integrity checking
- meta-data back-end
- network, rpc, fop, HA
- fom, reqh
- device io (stob)
- network raid repair
- security

containers

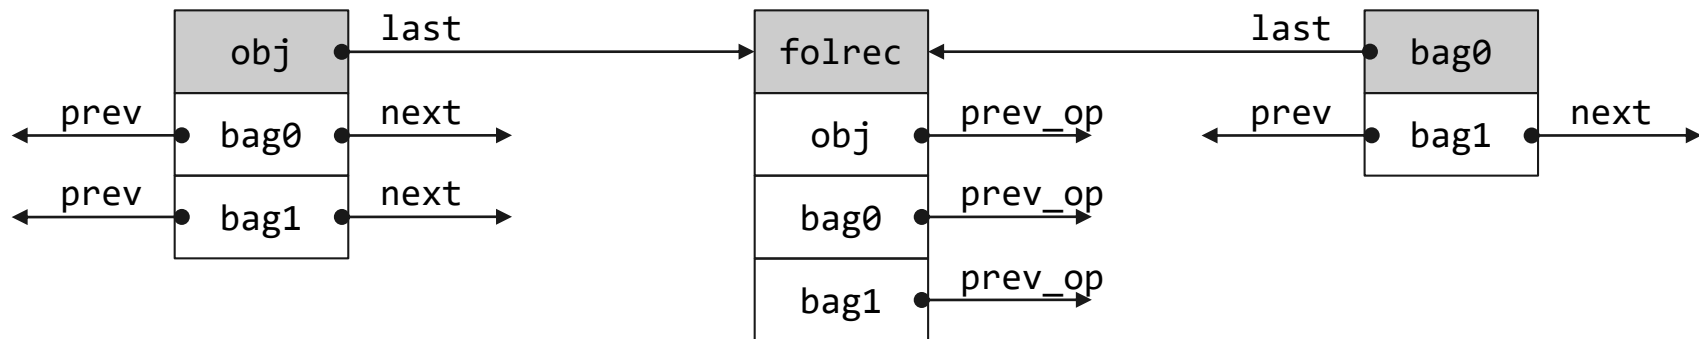
- container: an entity with a fid
- application fully controls containers
 - add an entity to a container;
 - remove an entity from a container;
 - list container elements;
 - execute bulk operation on container contents (compute-in-container)
- arbitrary "topology" (only restriction: no duplicates)

Non-properties:

- separate fid-space for contents
- strong isolation guarantees

containers implementation

- component containers, bag:
 - entities with fids,
 - lazily created when an entity is added to the container
- all local entities are linked together
- file operations log records (folrecs) for operations on container contents are linked



containers rationale

- efficiently identify containers to which an object belongs
- flexible container nesting
- concurrent mass-operations on containers
- container history is traceable
- fdmi filters on containers
- function shipping to a container

Don't have:

- ordering of container contents
- ordered enumeration
- "implicit" containers (e.g., "all objects on this server")

containers use case

- application creates a container to track "entities of interest"
- entities are added to the container
- fdmi plugin is registered to receive notifications of all updates to the container
- examples:
 - hsm: fileset of hot objects
 - replication: source fileset

components

- Motr client
- transactions (dtm)
- resource manager
- fdmi
- addb
- network raid, layouts
- containers
- **function shipping**
- lingua franca
- integrity checking
- meta-data back-end
- network, rpc, fop, HA
- fom, reqh
- device io (stob)
- network raid repair
- security

function shipping

- move computation closer to data (compute-in-storage)
- reduce network transmission overhead
- application structuring mechanism
- bulk computation on all Motr client entities:
 - object: function on data blocks
 - index: function on key-value records
 - container: function on member fids
- built-in fault tolerance

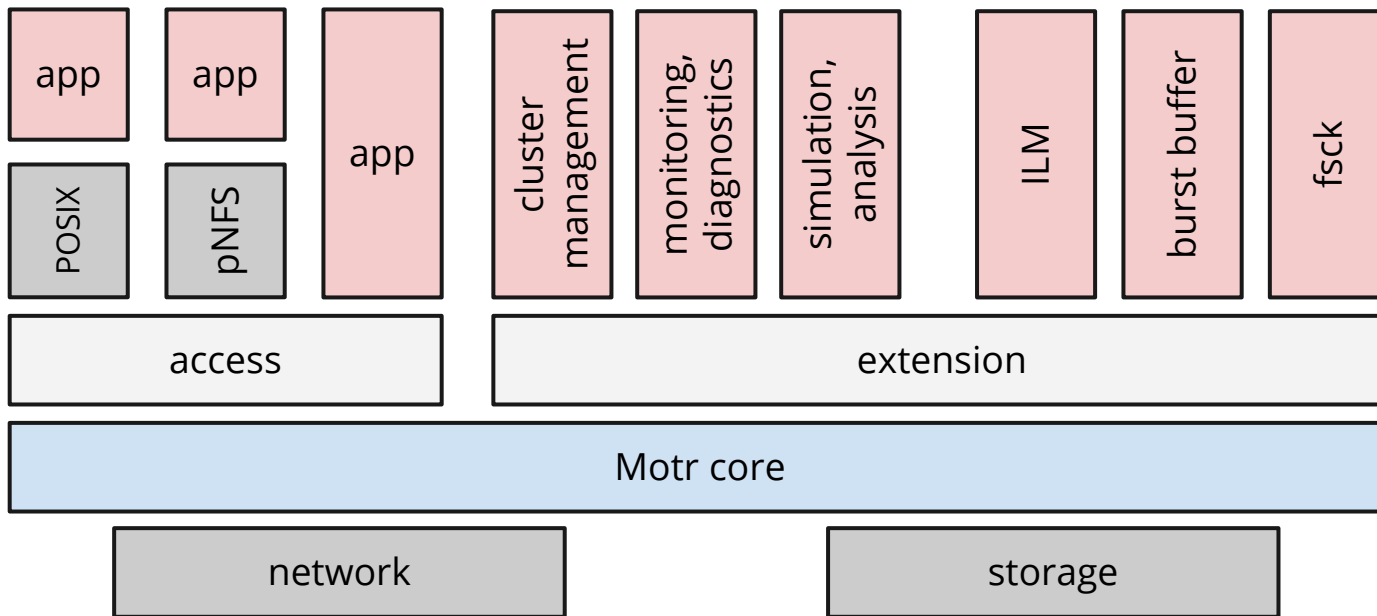
function shipping implementation

- computations are first class Motr client entities
 - unique fid, globally addressable
 - registered dynamically by an application
- low level trusted mechanism:
 - dynamically load shared library into Motr service process
 - invoke computations remotely, argument-result passing
- untrusted mechanism:
 - run untrusted code (e.g., Python) in a separate address space
- client uses layouts to start execution and recover from failures

components

- Motr client
- transactions (dtm)
- resource manager
- fdmi
- addb
- network raid, layouts
- containers
- function shipping
- **lingua franca**
- integrity checking
- meta-data back-end
- network, rpc, fop, HA
- fom, reqh
- device io (stob)
- network raid repair
- security

lingua franca



Multiple front-ends:
interoperability, common meta-data mechanism

lingua franca

- POSIX file system: *via* pNFS
 - straightforward semantics
 - concurrency:
 - concurrent reads
 - one writer
 - concurrent writers
- POSIX file system: *via* S3
 - interpret pathname as an URL (object key)
 - mapping of security and ownership attributes
- S3: *via* POSIX: parse URL as a pathname
- POSIX file system: *via* HDF5

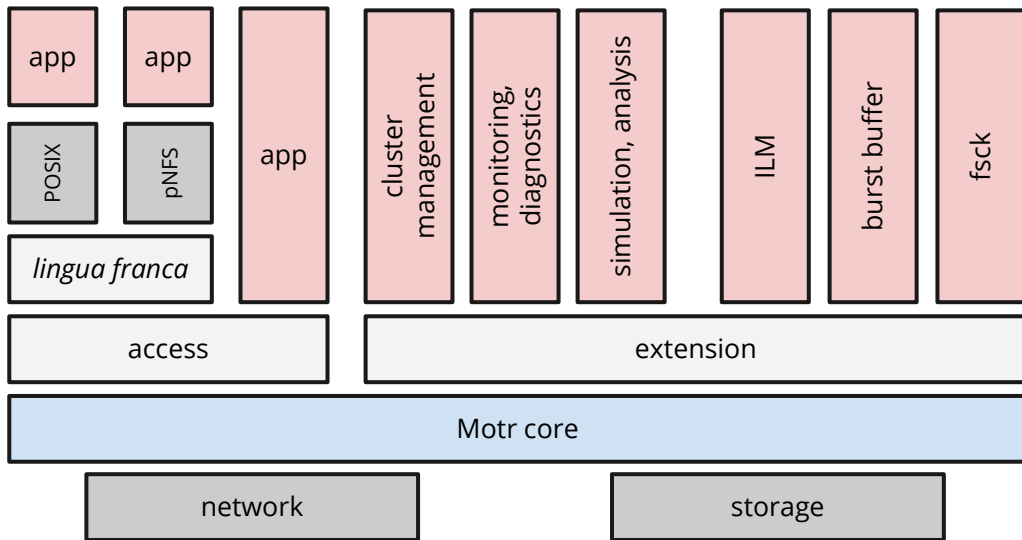
lingua franca

- S3 object: "blue_bucket/finance-9a84c723ee89f4b723b46cc5f1642b3"
- mount S3 store as POSIX

```
$ cd blue_bucket
$ ls -l finance-9a84c723ee8*

-rw----- 1 satoshi satoshi 285 January 03 2009 finance-9a84c723ee89f4b723b46cc5f1642b3
```
- enumerate and locate objects
- access attributes
 - system attributes: layout, containers;
 - common attributes: size;
 - front-end specific attributes: POSIX nlink
- efficiency: remote meta-data lookups are expensive

lingua franca



POSIX, pNFS, CIFS, S3/SWIFT, MPI Object IO, HDF5, MySQL, block device

lingua franca implementation

- what are the *entities* managed by an FE?
- how entities are named?
- how entities are organised: tree, graph, array?
- how attributes are associated with entities?
 - different FEs have different sets of native attributes
 - an FE wants to add attributes to foreign entities
 - an FE wants to interpret foreign attributes
 - some attributes are shared by multiple FEs
- core system has its own attributes

lingua franca implementation

- a domain of files (entities), a file identified by a 128-bit *fid*
- two indices:
 - name-space (ns): records file attributes
 - object index (oi): maps from the file fid to all its names

ns index structure:

key	value
FRONTENDID+FNAME+ATTRCLASS+ATTRID	attribute value

oi index structure:

FID+FRONTENDID+FNAMEID	FRONTENDID+FNAME
------------------------	------------------

lingua franca implementation: S3

- FRONTENDID: 3
- ATTRCLASS: 3
- FNAME: Object-URI: bucketname+NUL+object_key+NUL

object: "app_bucket/statement.xls", fid 0x60000:0x17ae76d0f

ns index:

key	value	comment
3app_bucket\0statement.xls\0!fid	0x60000:0x17ae76d0f	the fid of the object "statement.xls"
3app_bucket\0statement.xls\03content-length	3201526	object size in bytes
3app_bucket\0statement.xls\03content-md5	0x62ae2f12137738a9:0x173f5224f81446aa	md5 checksum
3app_bucket\0statement.xls\03last-modified	2017-04-26-15:29:30.85267	last modification timestamp
3app_bucket\0statement.xls\03...		other S3 attributes

lingua franca implementation: S3

object: "app_bucket/statement.xls", fid 0x60000:0x17ae76d0f

oi index:

key	value	comment
3:00000000000060000:0000000017ae76d0:30	3app_bucket\statement.xls	name of the object

lingua franca implementation: POSIX

- FRONTENDID: P
- ATTRCLASS: P
- FNAME: parent_directory_fid+name_in_the_directory

object: "/etc/passwd", fid 0x70000:0x322e1673fd

parent directory: "/etc", fid: 0x70000:0x18203a6485

ns index:

key	value	comment
P:70000:18203a6485:passwd!fid	0x70000:0x322e1673fd	the fid of the object "/etc/passwd"
P:70000:18203a6485:passwdPsize	16523	file size
P:70000:18203a6485:passwdPatime	2017-04-26-15:29:30.85267	access time
P:70000:18203a6485:passwdP...		other POSIX attributes

lingua franca implementation: POSIX

object: "/etc/passwd", fid 0x70000:0x322e1673fd

oi index:

key	value	comment
P:00000000000070000:0x000000322e1673fd:P0	P:70000:18203a6485:passwd	name of the object
P:00000000000070000:0x000000322e1673fd:P1	P:70000:18203a6485:passwd.1	another name: hard-link

```
# ln /etc/passwd /etc/passwd.1
```

lingua franca implementation: features

- each attribute is a separate key-value pair: flexibility
- new attributes can be added to existing files
 - without breaking compatibility
- attributes can be enumerated (NEXT operation)
 - an FE can selectively handle attributes it understands
- contents of a directory can be enumerated
- new names can be added to a file

components

- Motr client
- transactions (dtm)
- resource manager
- fdmi
- addb
- network raid, layouts
- containers
- function shipping
- lingua franca
- **integrity checking**
- meta-data back-end
- network, rpc, fop, HA
- fom, reqh
- device io (stob)
- network raid repair
- security

integrity checking

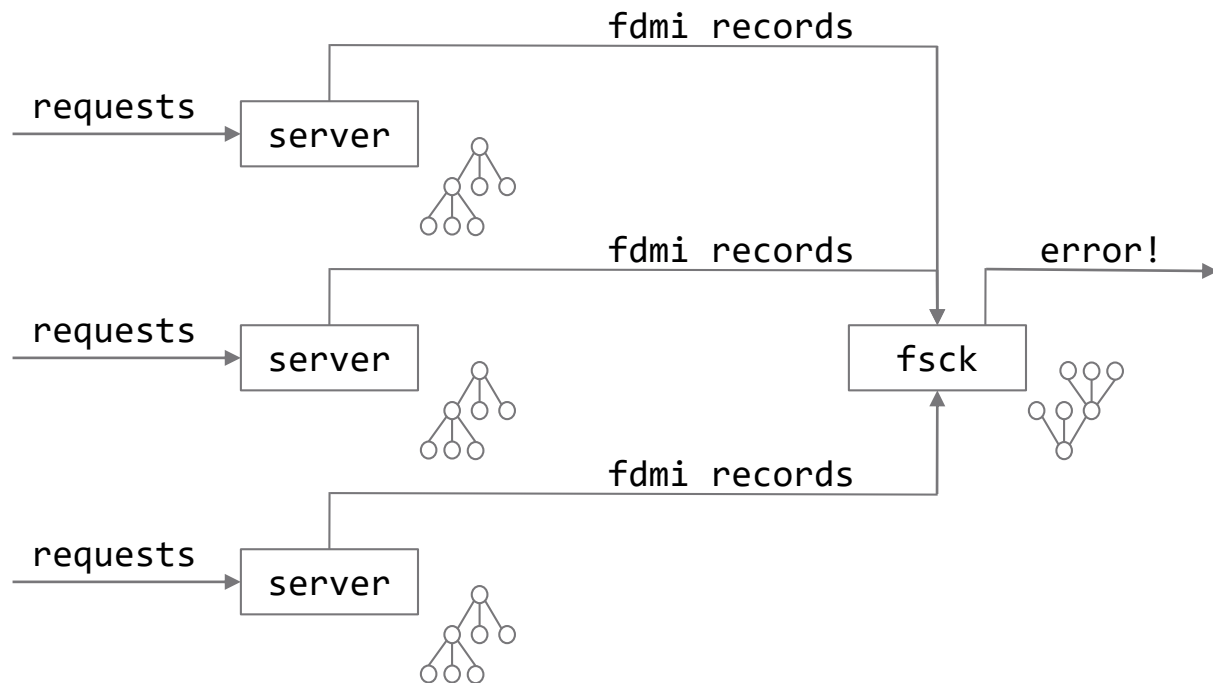
- redundancy, fancy metadata: not an answer (has been tried)
 - bugs (more important over time)
 - recovery from catastrophic failures
- traditional fsck
 - not distributed
 - specific to particular meta-data format
 - does not scale
 - time
 - space

integrity checking

- need scalable integrity checking
- run it all the time
- on dedicated separate nodes (horizontal scalability)
 - maintain redundant "inverse" meta-data
 - update meta-data to match system evolution (fdmi)
 - detect inconsistencies
 - report, recover from redundancy
 - recover catastrophic failures
- usual redundancy: parity, checksums, background scrub
- fdmi: transactional coherence with the main state

- parallel programming

integrity checking



inverse meta-data

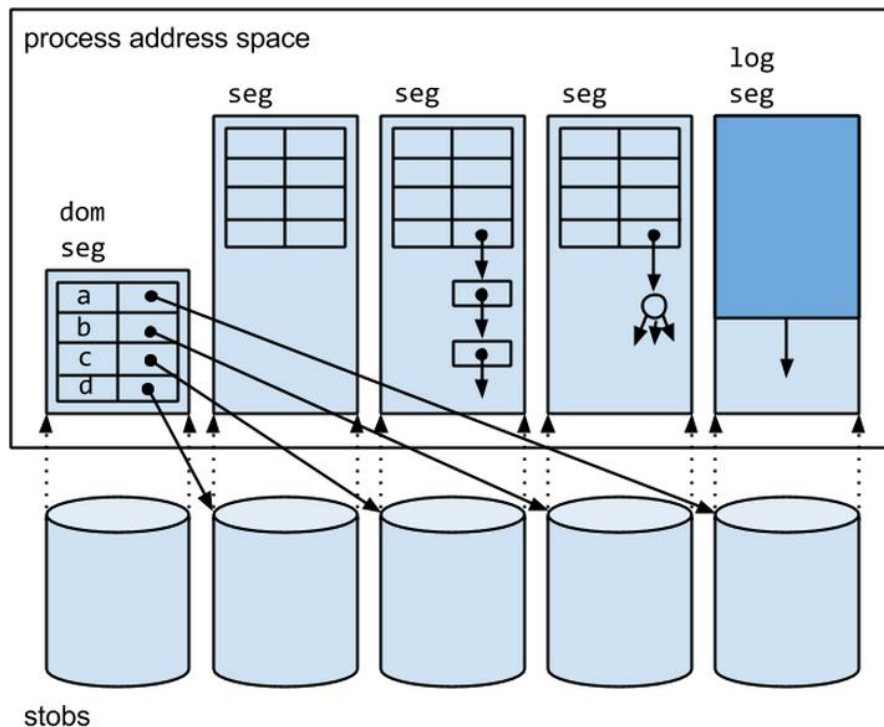
- block allocation
- pdclust structure
- key distribution
- b-tree structure
- application specific invariants
 - POSIX tree
 - hdf5

components

- Motr client
- transactions (dtm)
- resource manager
- fdmi
- addb
- network raid, layouts
- containers
- function shipping
- lingua franca
- integrity checking
- **meta-data back-end**
- network, rpc, fop, HA
- fom, reqh
- device io (stob)
- network raid repair
- security

backend

- segment
- transaction
- domain
- container
- WAL
- redo-only
- allocator
- btree



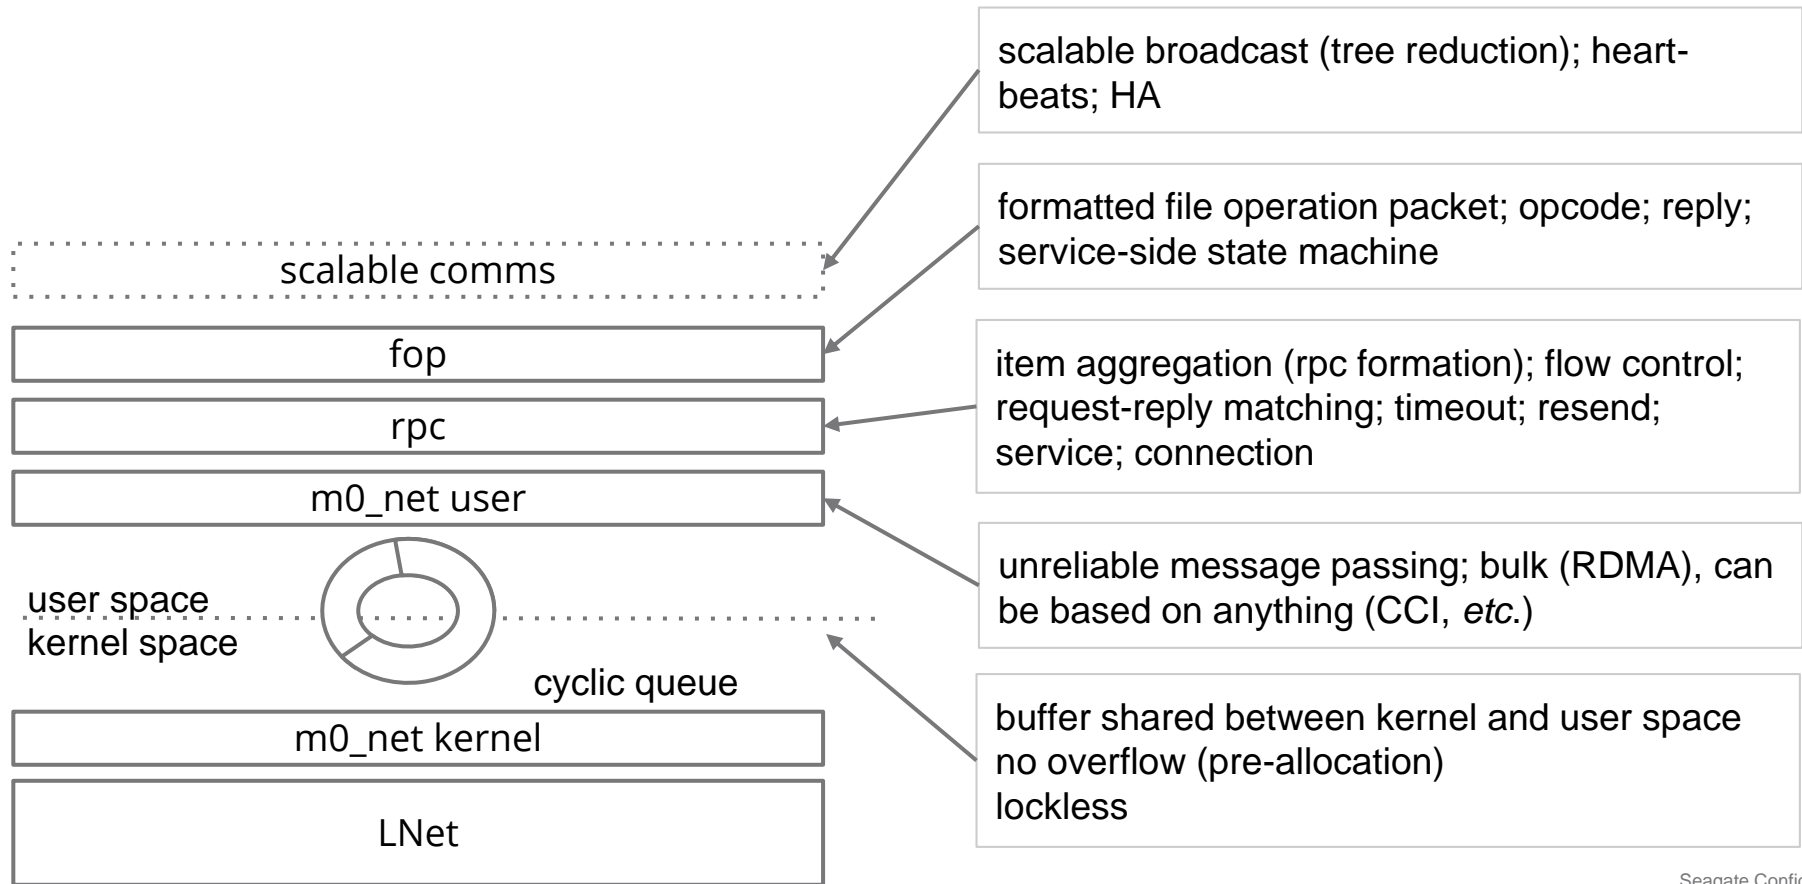
components

- Motr client
- transactions (dtm)
- resource manager
- fdmi
- addb
- network raid, layouts
- containers
- function shipping
- lingua franca
- integrity checking
- meta-data back-end
- **network, rpc, fop, HA**
- fom, reqh
- device io (stob)
- network raid repair
- security

comm

- network: LNet, 0-copy, unreliable message passing
- rpc
 - message packing (formation),
 - request-reply semantics
 - retransmit
- xcode: serialisation library
- fop: operation packet
- *reduce-broadcast*
 - communication and aggregation tree
 - provided by HA

comm



comm: xcode

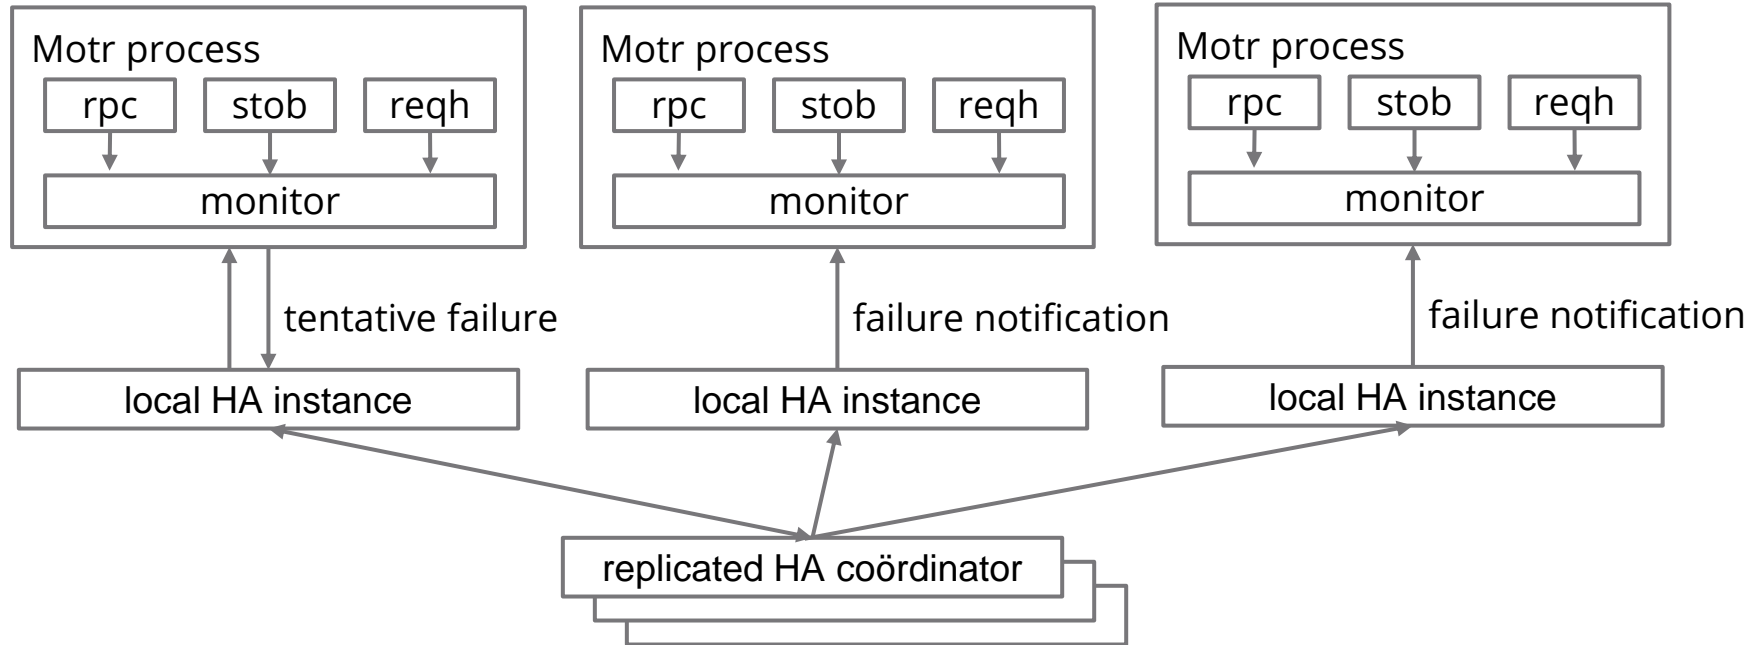
- a library adding a (modest) amount of reflection to C
- annotate header file foo.h
- gccxml: C parser
- generate foo_xc.h, foo_xc.c

```
struct m0_foo {  
    uint32_t      f_nr;  
    struct m0_bar *f_bar;  
} M0_XCA_RECORD;
```

```
static struct m0_xcode_type m0_foo_xc = {  
    .xct_name      = "m0_foo",  
    .xct_sizeof    = sizeof (struct m0_foo),  
    .xct_children  = {  
        [0] = { "f_nr",  &M0_XT_U32,  offsetof(struct m0_foo, f_nr) },  
        [1] = { "f_bar", &m0_bar_xc,  offsetof(struct m0_foo, f_bar) }  
    };  
};
```

- m0_xcode_{en,de}code(), m0_xcode_{find,print,read}()

comm: HA

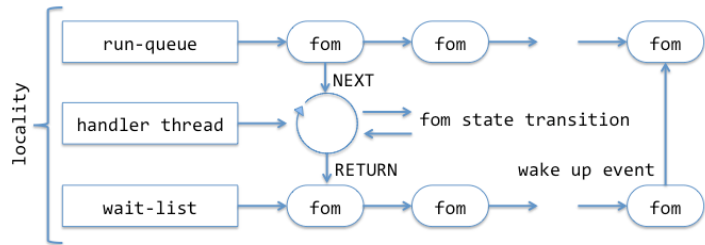
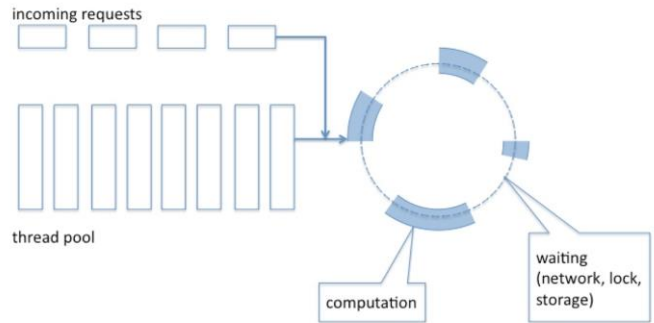


components

- Motr client
- transactions (dtm)
- resource manager
- fdmi
- addb
- network raid, layouts
- containers
- function shipping
- lingua franca
- integrity checking
- meta-data back-end
- network, rpc, fop, HA
- **fom, reqh**
- device io (stob)
- network raid repair
- security

fom

- thread-per-request:
 - multiple cores, NUMA,
 - locking,
 - cache ping-pong,
 - c10K, many threads
- reqh:
 - thread per core
 - non-blocking scheduler
 - locality of reference
 - load balancing
 - long-term scheduling



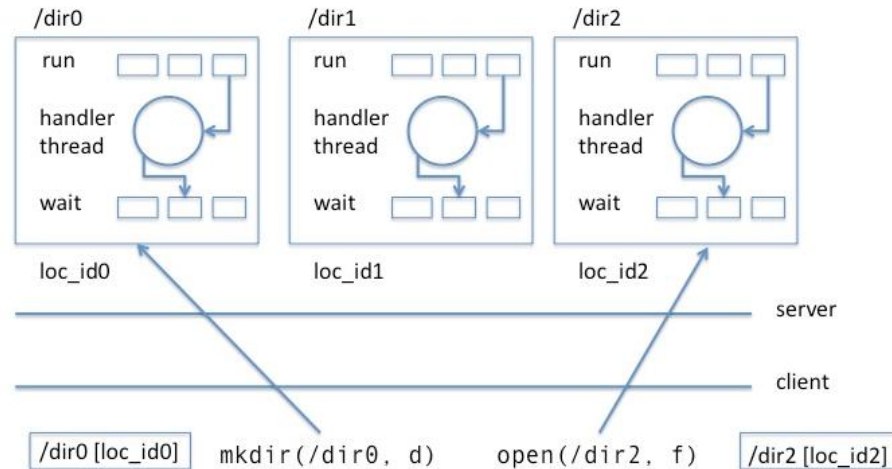
fom: request handler

```
while (!locality->shutdown) {
    while (have_work) {
        network_drain();
        stob_drain();
        fom_transition();
    }
    wait_for_work();
}

fom_transition(locality *l) {
    for_each(fom, l->run_queue) {
        good = goodness(l, fom);
        if (good > max_goodness) {
            best = fom;
            max_goodness =
                good;
        }
    }
    best->state(); /* state transition */
}

goodness(locality *l, fom *f) {
    if (abs(fom->next_block_nr, l->elevator_pos) < threshold)
        result += 1;
    if (l->pending_rpc[fom->next_nid] > 0)
        result += 1;
    result += take_deadline_into_account(l, f);
    result += take_priority_into_account(l, f);
    ...
}
```

fom: load balancing



- each locality is a small server
- clients talk to particular locality, using opaque identifier
- change identifier for load-balancing

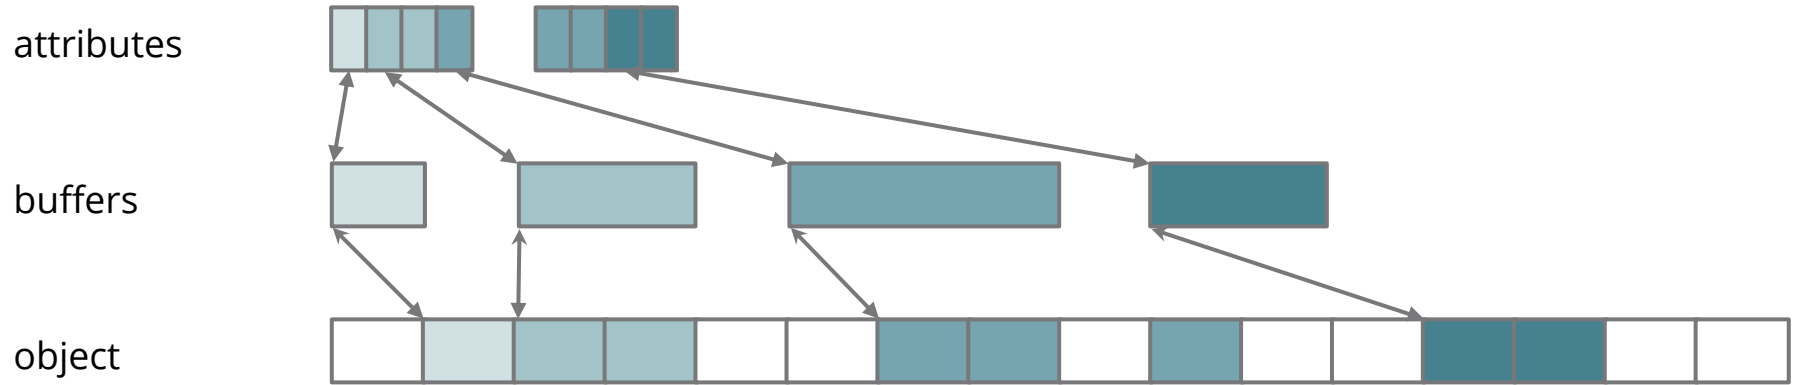
components

- Motr client
- transactions (dtm)
- resource manager
- fdmi
- addb
- network raid, layouts
- containers
- function shipping
- lingua franca
- integrity checking
- meta-data back-end
- network, rpc, fop, HA
- fom, reqh
- **device io (stob)**
- network raid repair
- security

stob

- array of data-blocks, $[0, 2^{64})$, initially a hole
- create, delete, read, write, alloc, free operations
- IO at block granularity
- no usual meta-data (attributes, *etc.*)
- block attributes can be used for checksums, encryption keys, hash fingerprints
- scatter-gather-scatter operations: data and block attributes

stob: IO



stob: implementations

- linuxstob (*aka* devstob)
 - stob = file
 - aio
- adstob (allocation data stob)
 - multiple stobs stored in a backend stob
 - block allocator balloc
 - based on ext4 mballoc

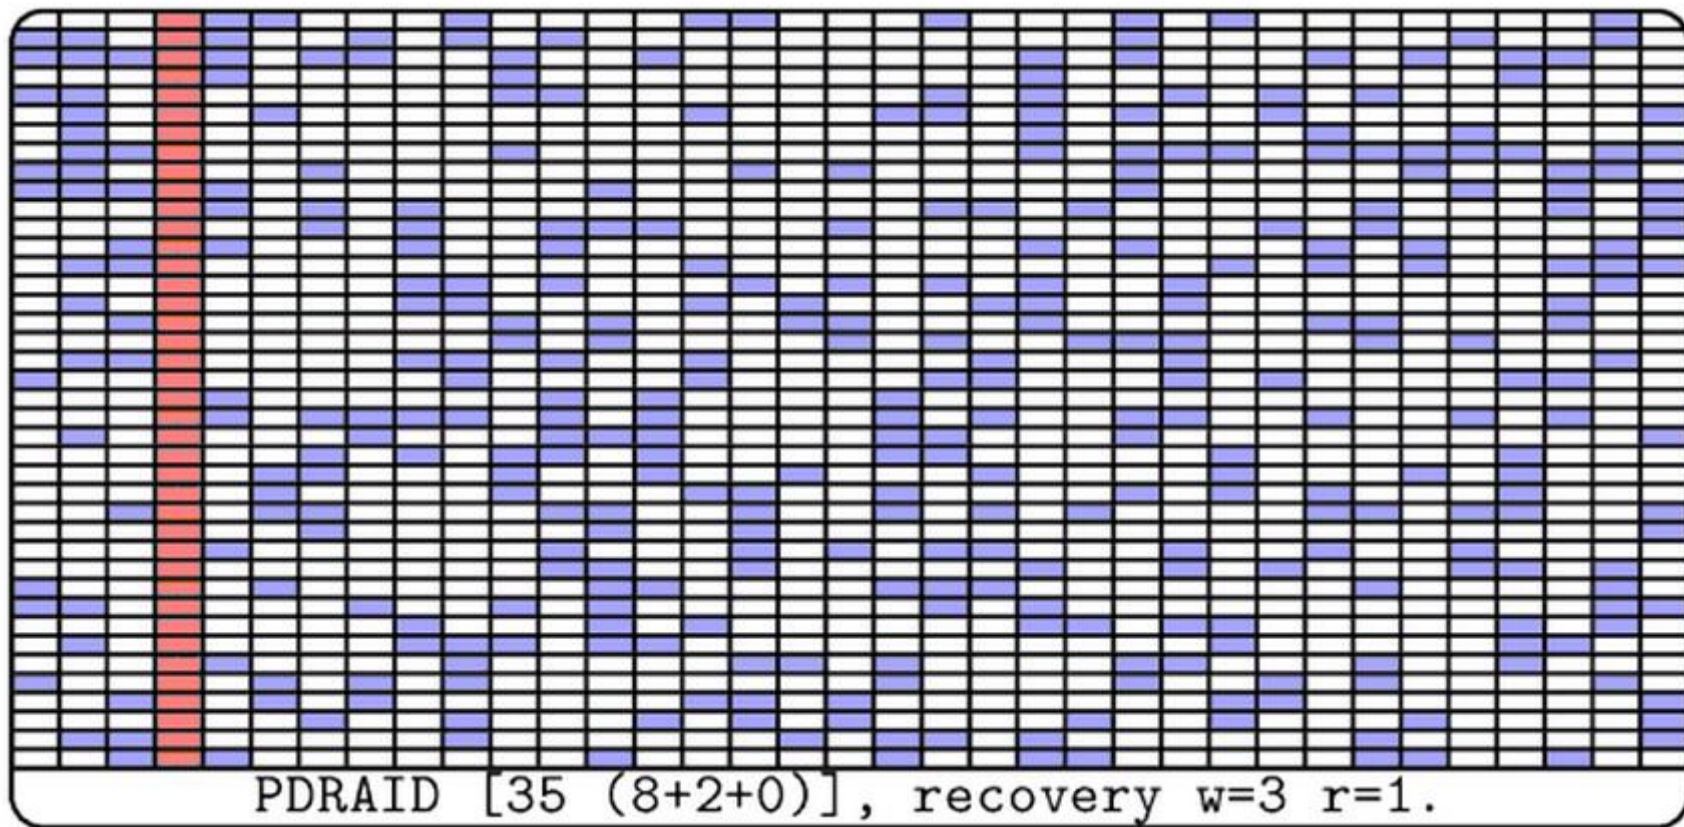
components

- Motr client
- transactions (dtm)
- resource manager
- fdmi
- addb
- network raid, layouts
- containers
- function shipping
- lingua franca
- integrity checking
- meta-data back-end
- network, rpc, fop, HA
- fom, reqh
- device io (stob)
- **network raid repair**
- security

sns

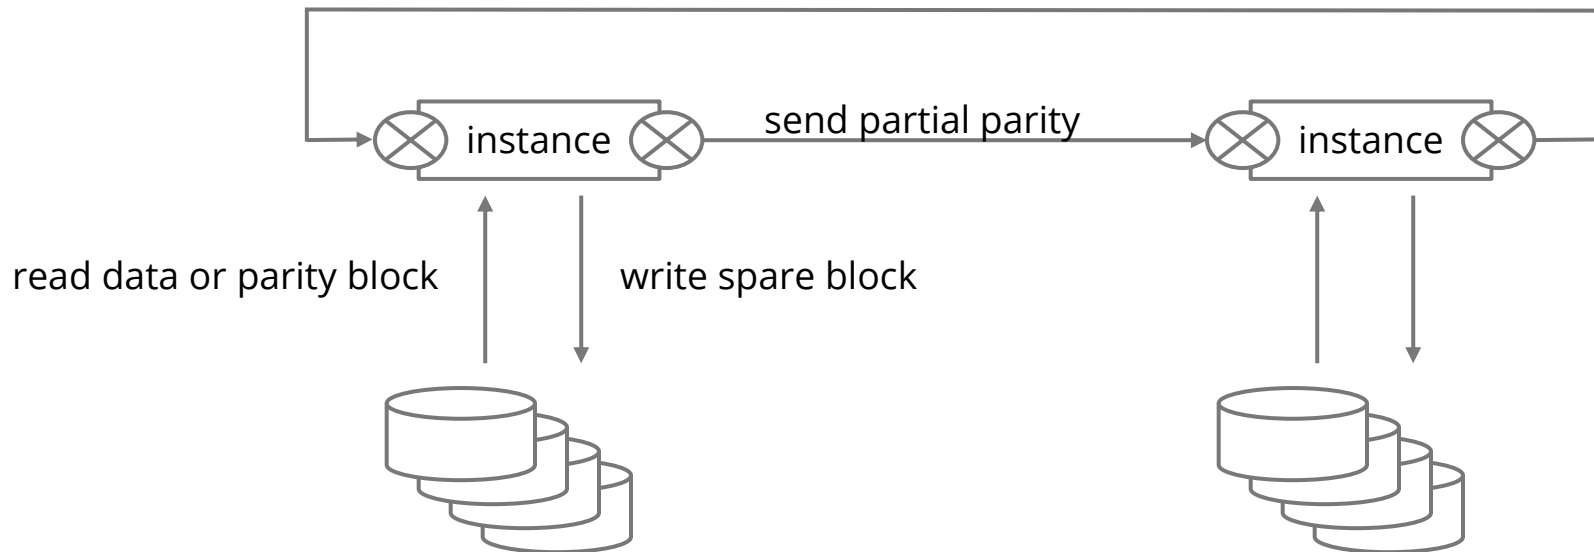
- guaranteed IO performance during repair
- fast repair
- copy machine
- repair
- rebalance
- pool
- *flattening*

sns: repair

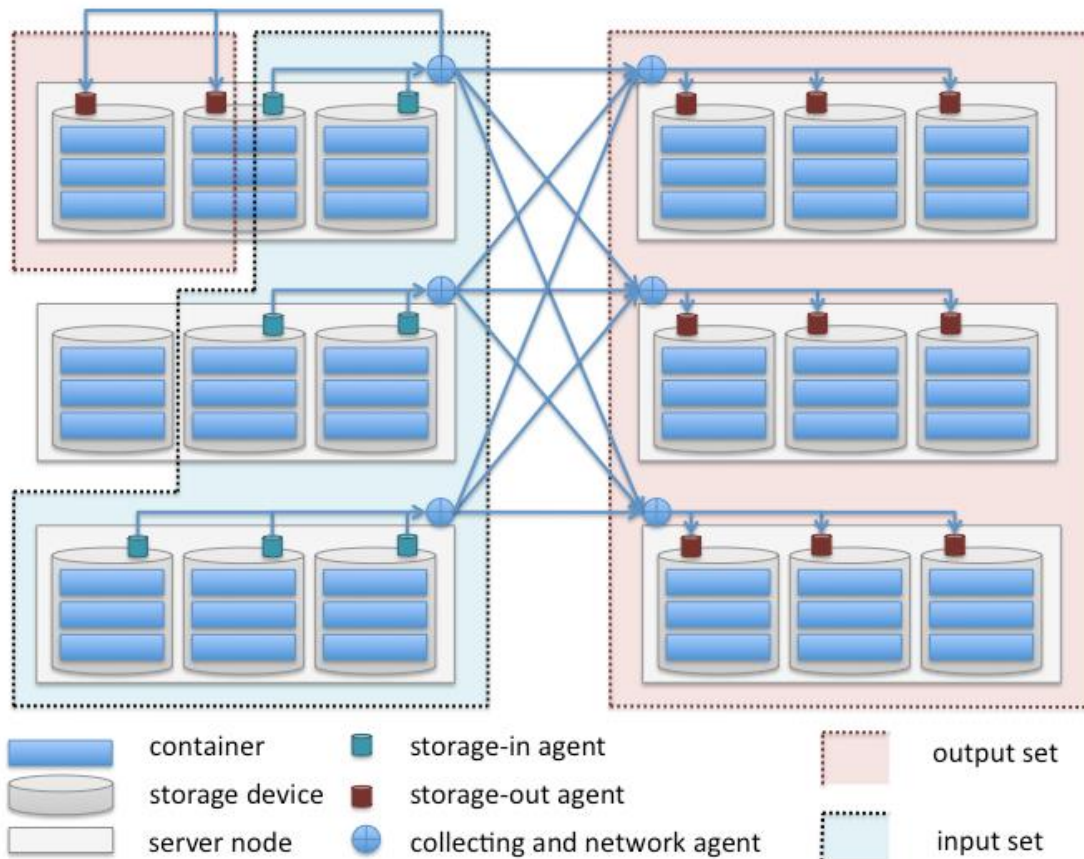


sns: copy machine

parity group

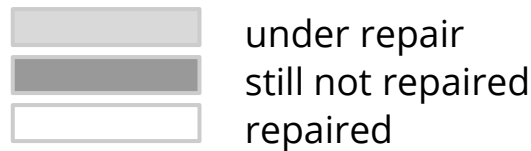
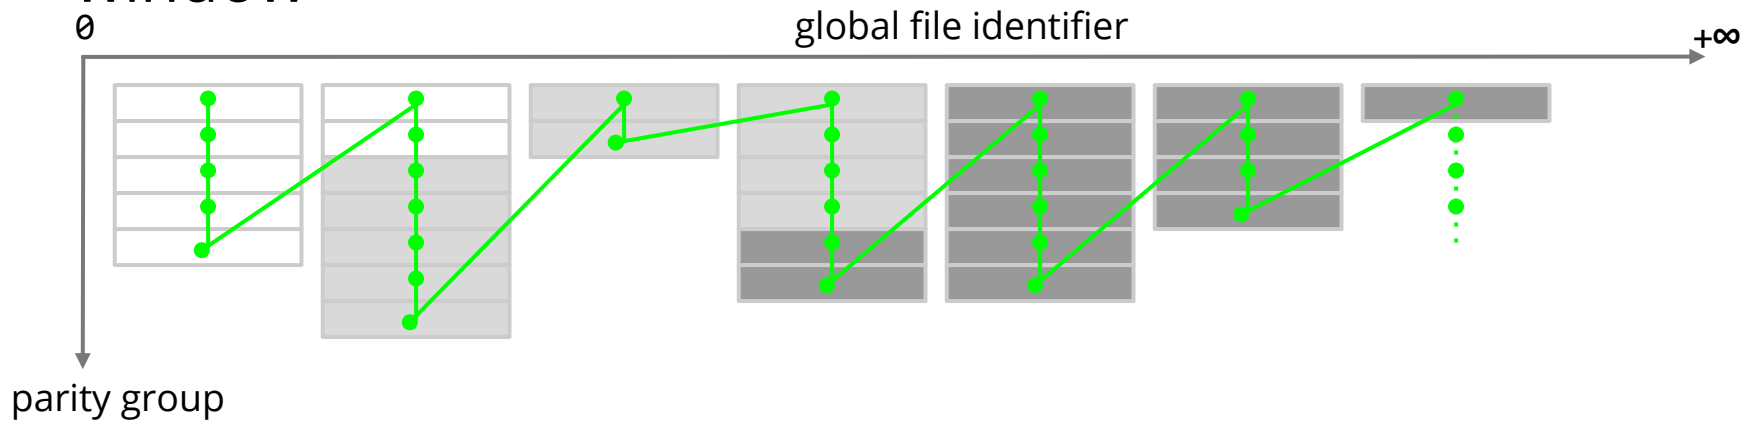


sns: network

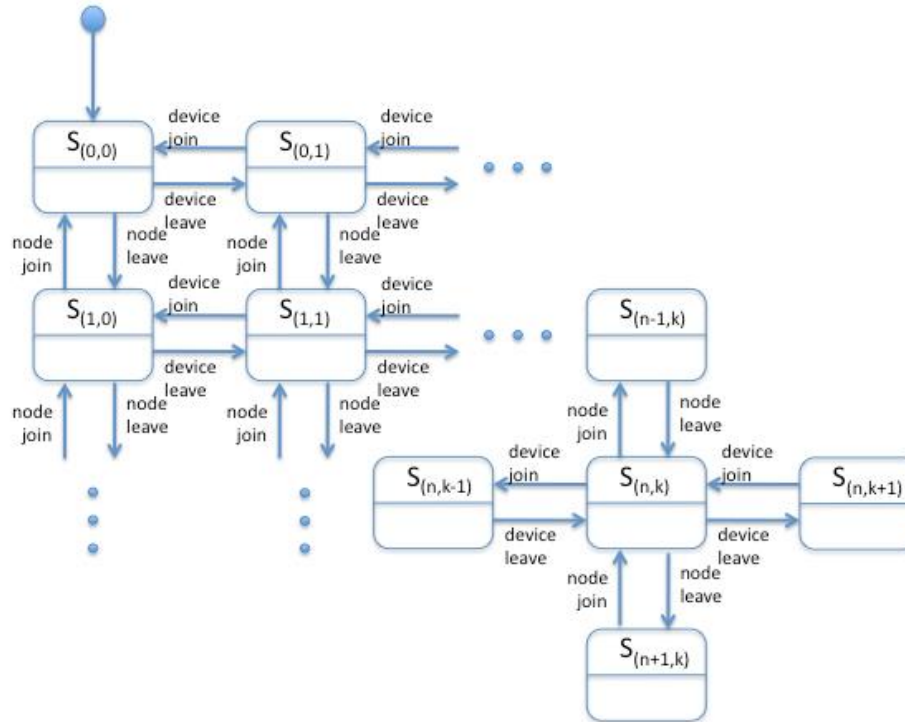


sns: coördination

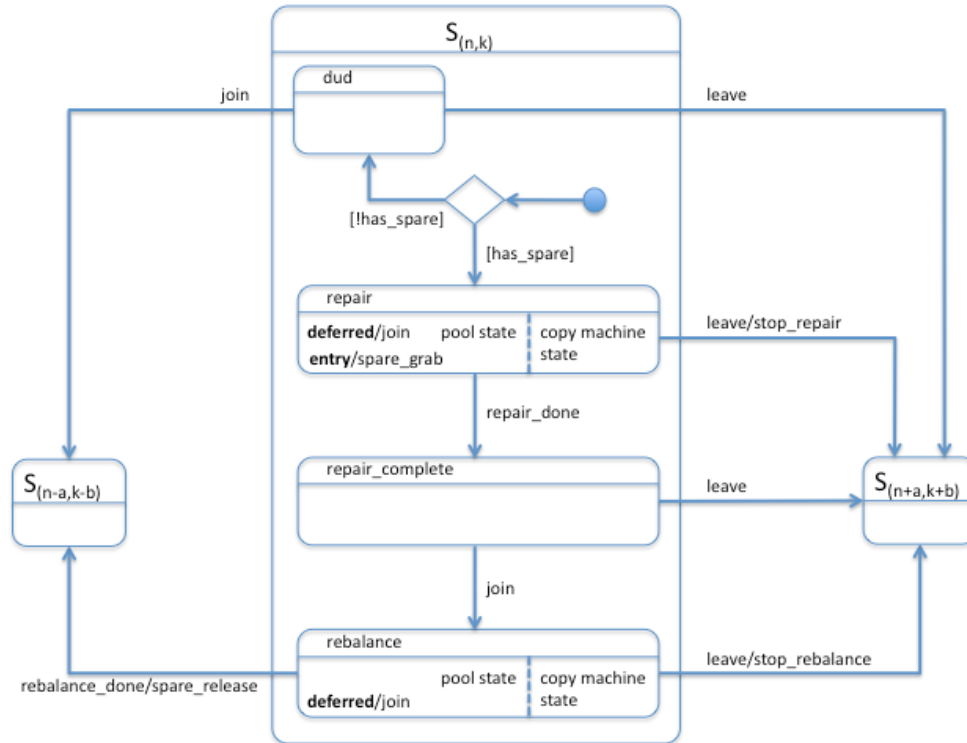
copy machine sliding
window



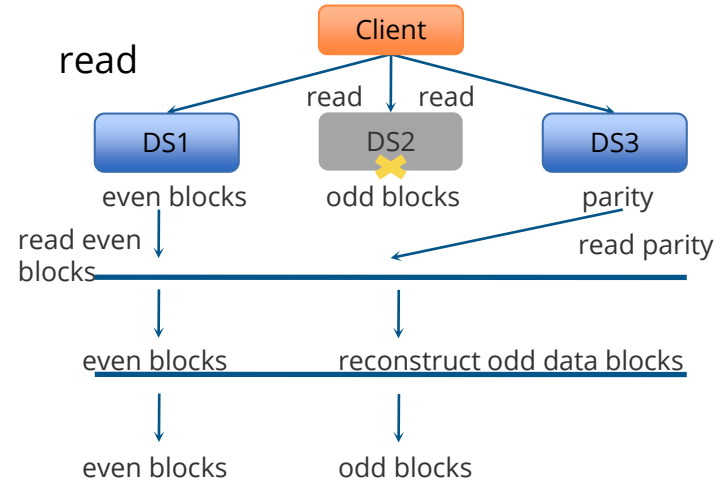
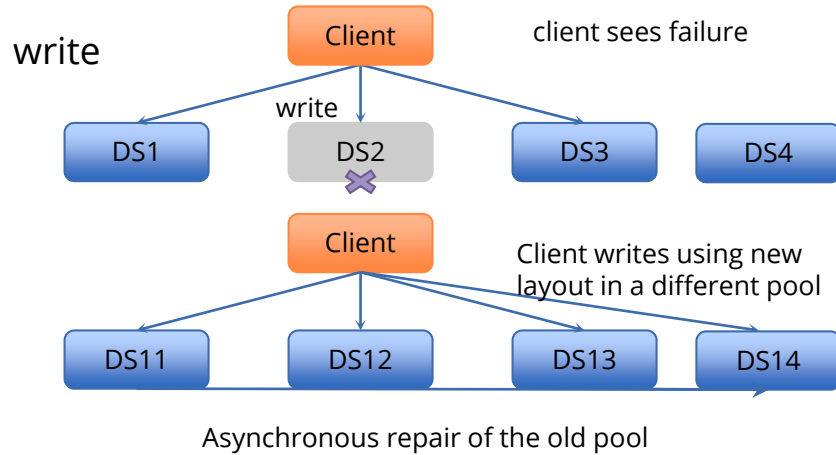
sns: failure state machine



sns: failure in detail



sns: nba

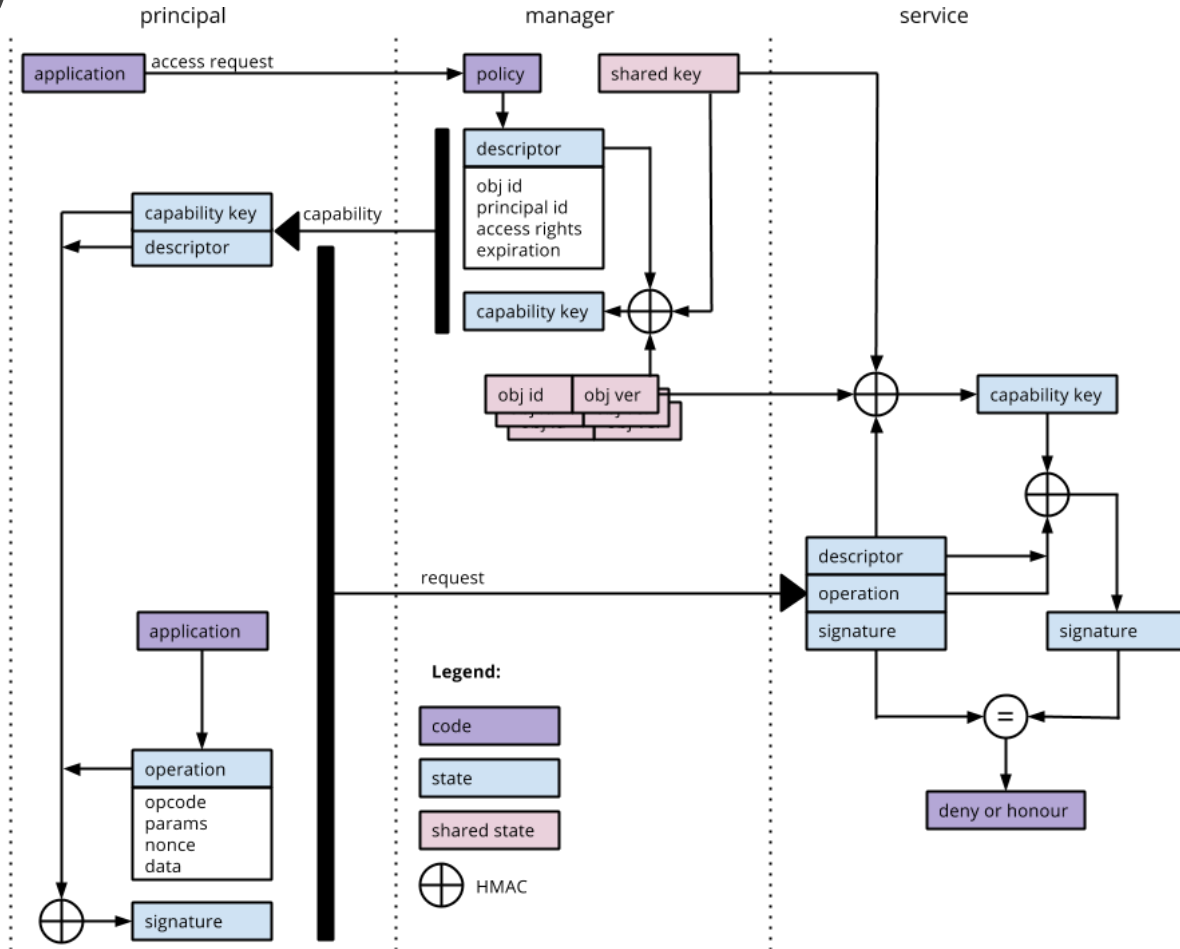


- client senses failure by timeout, notifies HA
- how new layout is selected? Layout formula
- composite layout with list of extents
- flattening: restore simple layout

components

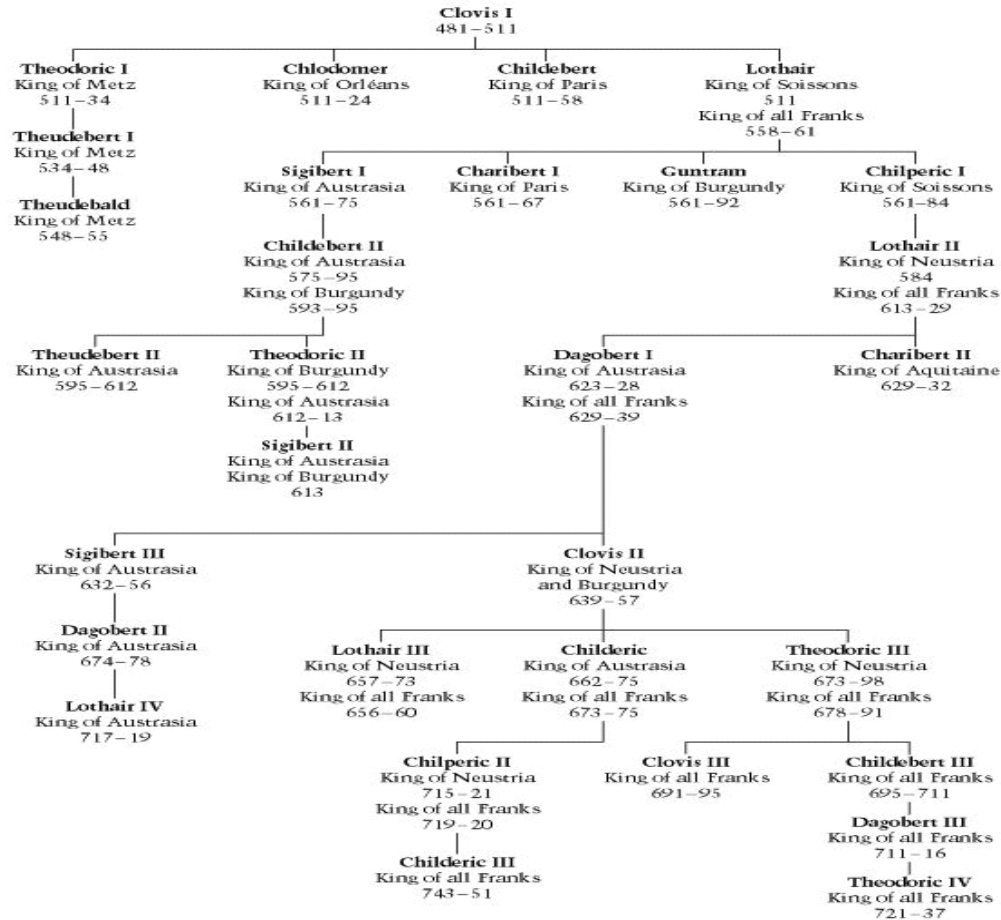
- Motr client
- transactions (dtm)
- resource manager
- fdmi
- addb
- network raid, layouts
- containers
- function shipping
- lingua franca
- integrity checking
- meta-data back-end
- network, rpc, fop, HA
- fom, reqh
- device io (stob)
- network raid repair
- **security**

security



questions?

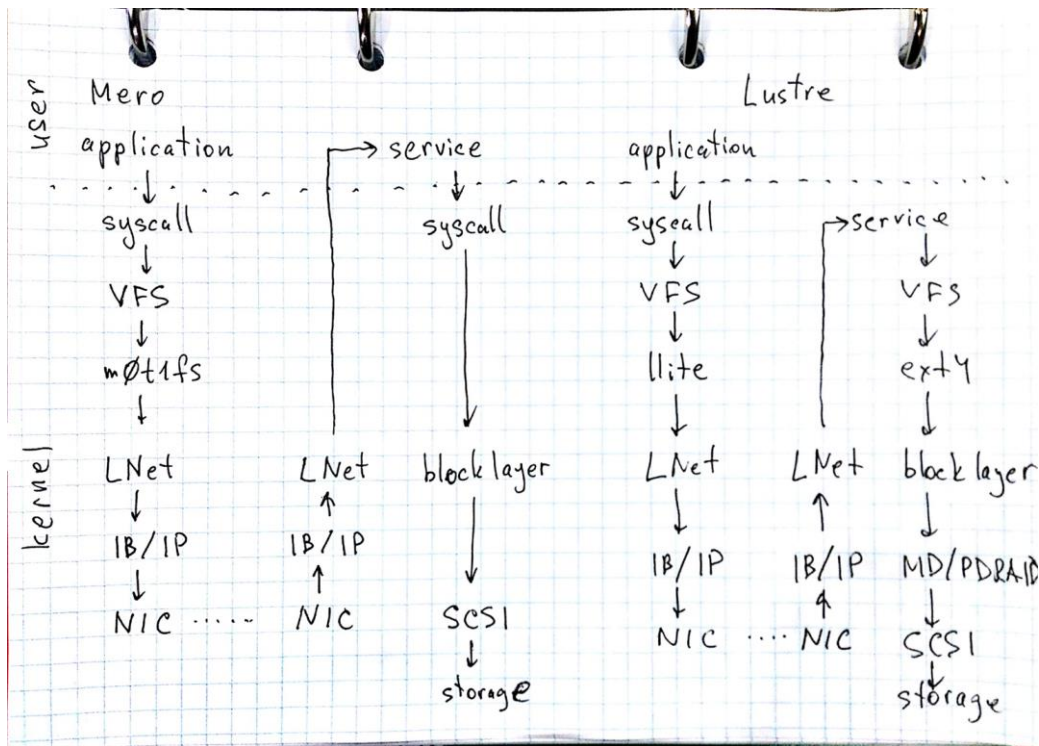
THE MEROVINGIAN KINGS



use cases

- IO data-flow end-to-end
- lookup path through all layers
- LOMO/WOMO: server and client failure

data-flow



lookup: m0t1fs

```
$ ls -l foo
```

- stat("foo", &stbuf);
 - m0t1fs_lookup(dir, dentry, nd)
 - fop = m0__fop_alloc(&m0_fop_lookup_fopt);
 - m0_rpc_post(fop)
 - m0_rpc_item_wait_for_reply(fop)
 - m0t1fs_iget()

lookup: rpc out

- `m0_rpc_post(fop)`
 - `m0_rpc_item_send()`
 - `m0_rpc_item_start_timer()`
 - `m0_rpc_frm_enq_item()`
 - `frm_insert()`
 - `queue add [URGENT, BOUND]`
 - `frm_balance()`
 - `if (ready) frm_fill_packet()`
 - `m0_rpc_packet_add_item()`
 - `frm_packet_ready()`
 - `m0_net_buffer_add()`

lookup: rpc in

- `buf_recv_cb()` (= `[M0_NET_QT_MSG_RECV]`)
 - `net_buf_received()`
 - `packet_received()`
 - `item_received()`
 - `m0_rpc_slot_reply_received()`
 - `item_find(item_xid(reply))`
 - `req->rio_replied()`
 - `m0_rpc_item_change_state(req, M0_RPC_ITEM_REPLIED)`
 - `m0_rpc_item_wait_for_reply()`

lookup: network out

- `m0_net_buffer_add()`
 - `m0_net_tm_tlink_init_at_tail(buf, ql)`
 - `nx_ops->xo_buf_add(buf) (= nlx_xo_buf_add)`
 - `M0_NET_QT_MSG_SEND`
 - `nlx_core_buf_msg_send()`
 - `nlx_ucore_ioctl(M0_LNET_BUF_MSG_SEND)`

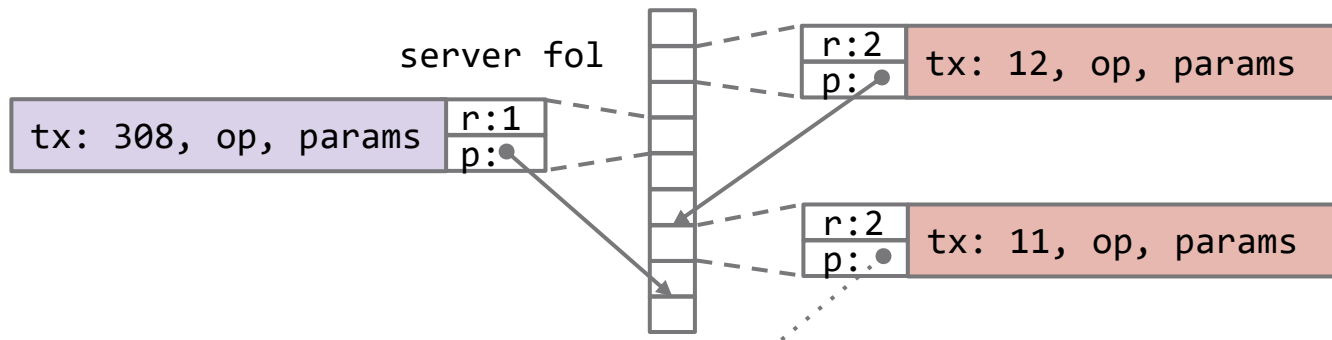
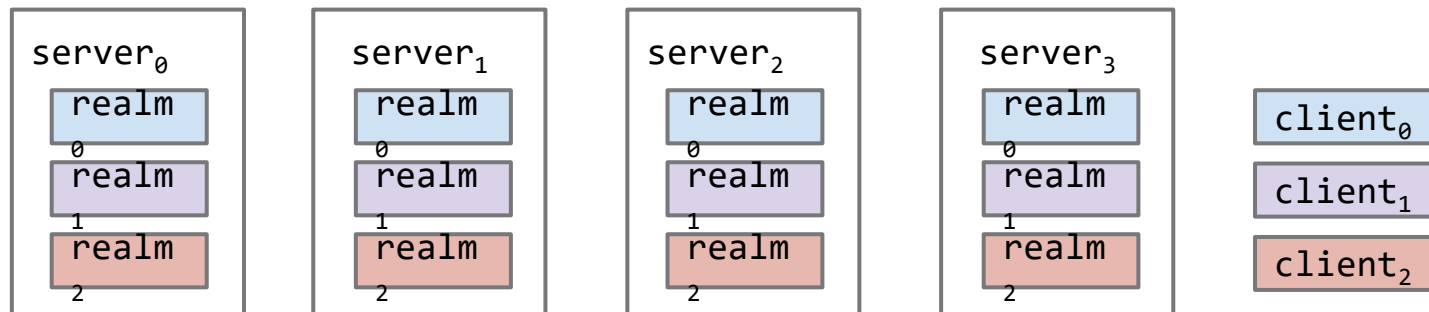
lookup: network in

- `nix_tm_ev_worker()`
 - `!queue.empty()` or `semaphore_down()`
 - `queue get`
 - `m0_net_tm_event_post()`
 - `buf_rcv_cb()` (= `[M0_NET_QT_MSG_RECV]`)

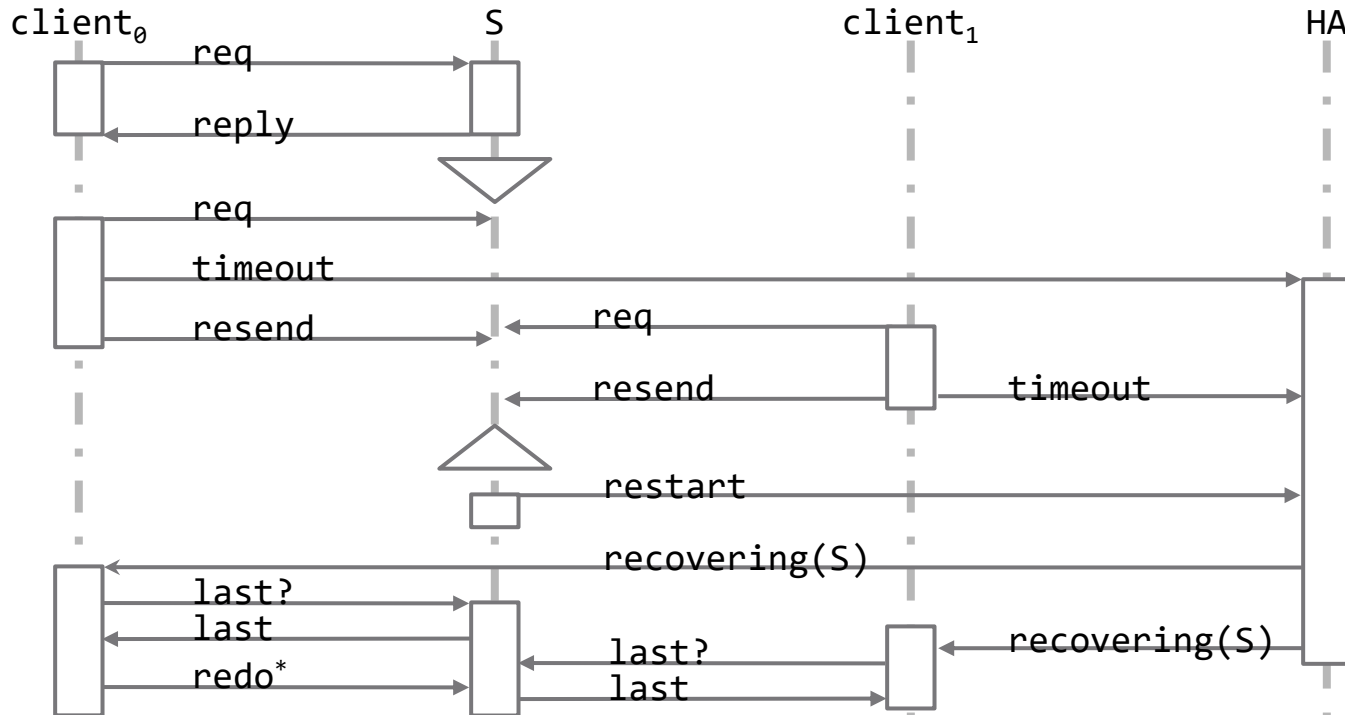
lookup: service side

- `m0_rpc_item_dispatch()`
 - `m0_reqh_fop_handle()`
 - `fop->fto_create(fop, &fom)`
 - `m0_fom_queue(fom)`
- `loc_handler_thread()`
 - `fom_exec(fom)`
 - `fom->fo_ops->fo_tick(fom)`
 - `m0_md_tick_getattr()`
 - `m0_fom_tick_generic()`
 - `m0_cob_locate()`
 - `m0_be_btree_lookup()`

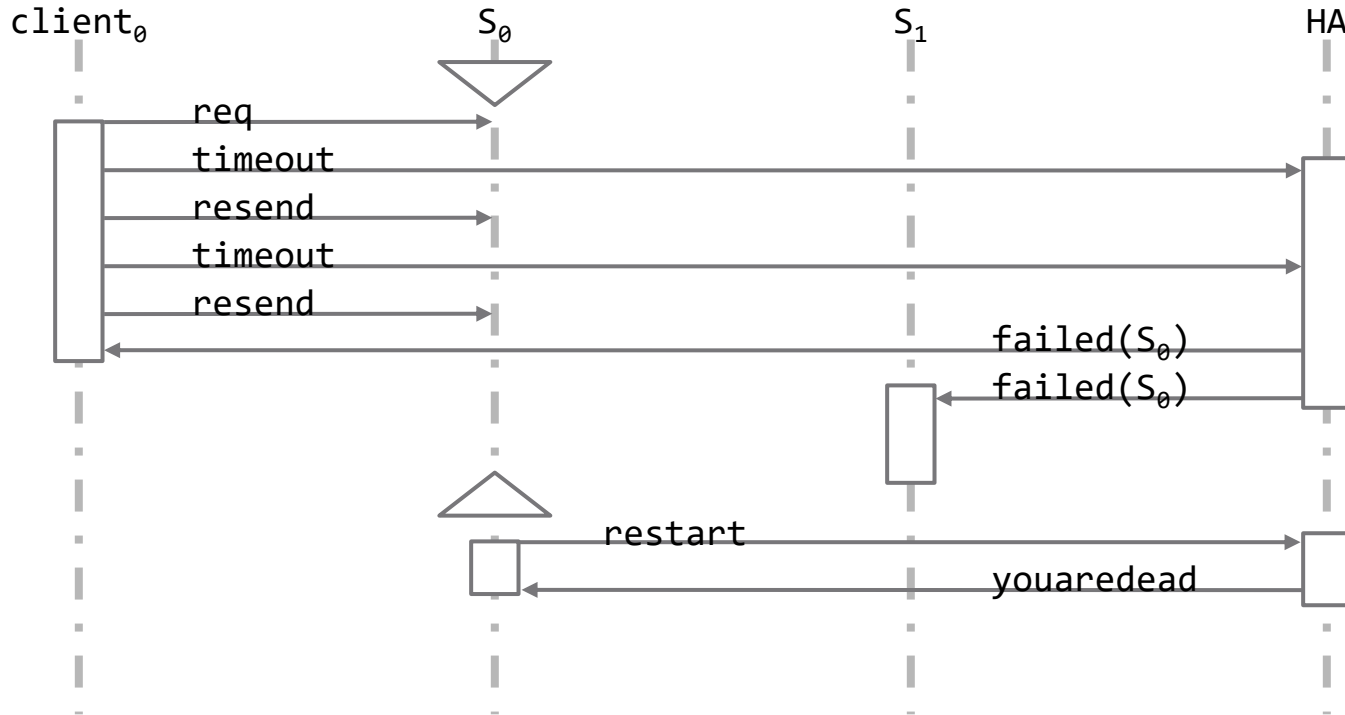
LOMO/WOMO realms



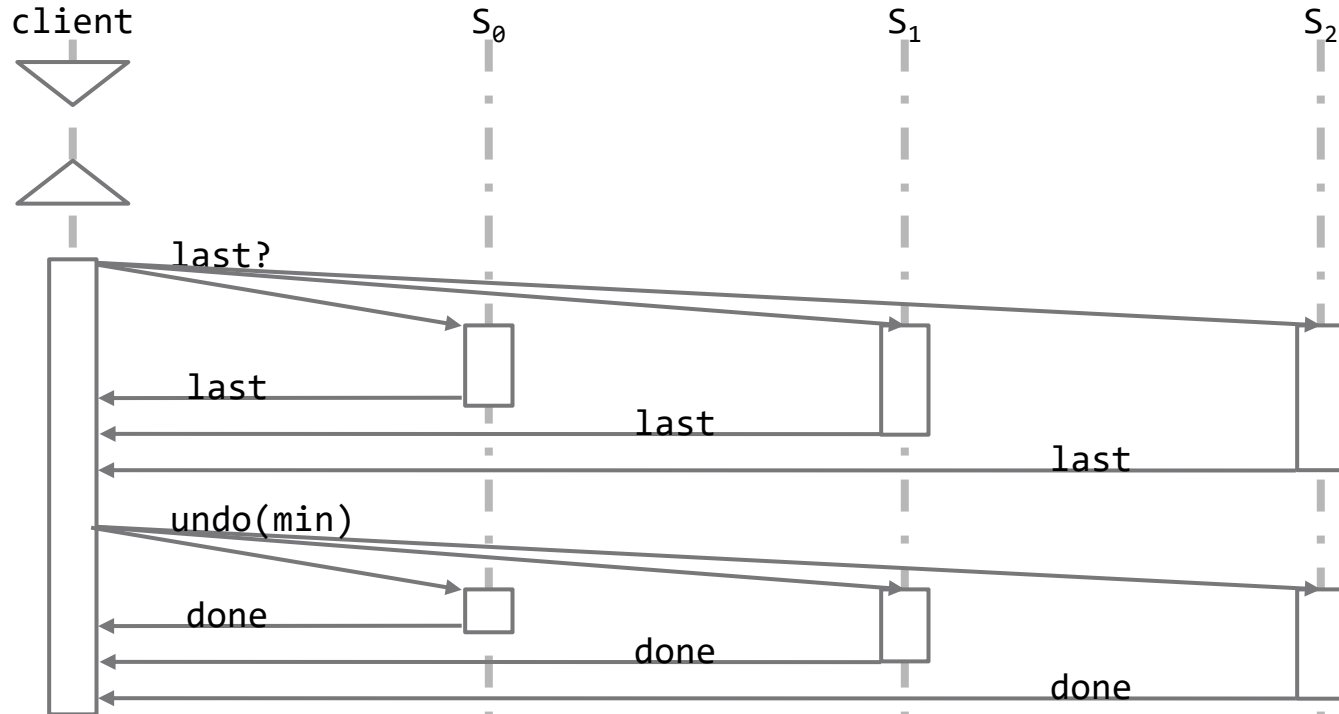
transient server failure



permanent server failure



client failure



io:

